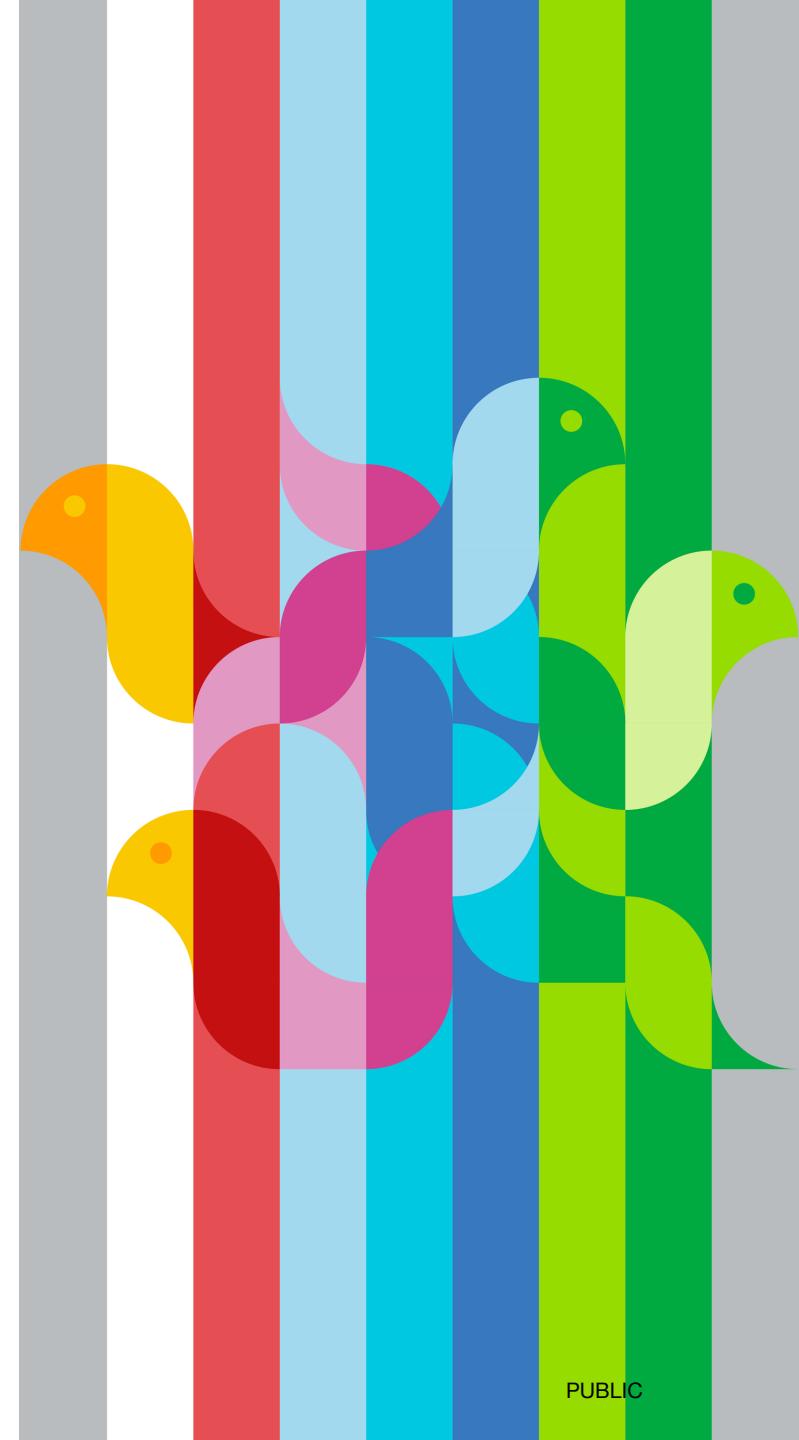


Rustifying Python: A Practical Guide to Achieving High Performance While Maintaining Observability

Max Höhl
SAP SE





/whme/PyConDE-2025

PUBLIC

“

**Don't force me to
deal with your shiny
language of the day.
Maintaining multi-
language projects is
a pain**

- Christoph Hellwig



/whme/PyConDE-2025

PUBLIC

- How to identify parts suitable for Rust



/whme/PyConDE-2025

PUBLIC

- How to identify parts suitable for Rust
- How to migrate to Rust and deploy to production



/whme/PyConDE-2025

PUBLIC

- How to identify parts suitable for Rust
- How to migrate to Rust and deploy to production
- Lessons learned



/whme/PyConDE-2025

PUBLIC









1 – Conflict

2 – Cluster Utilization – 3 4 5 6 7 8



Cluster Utilization



Data management

1 – Conflict

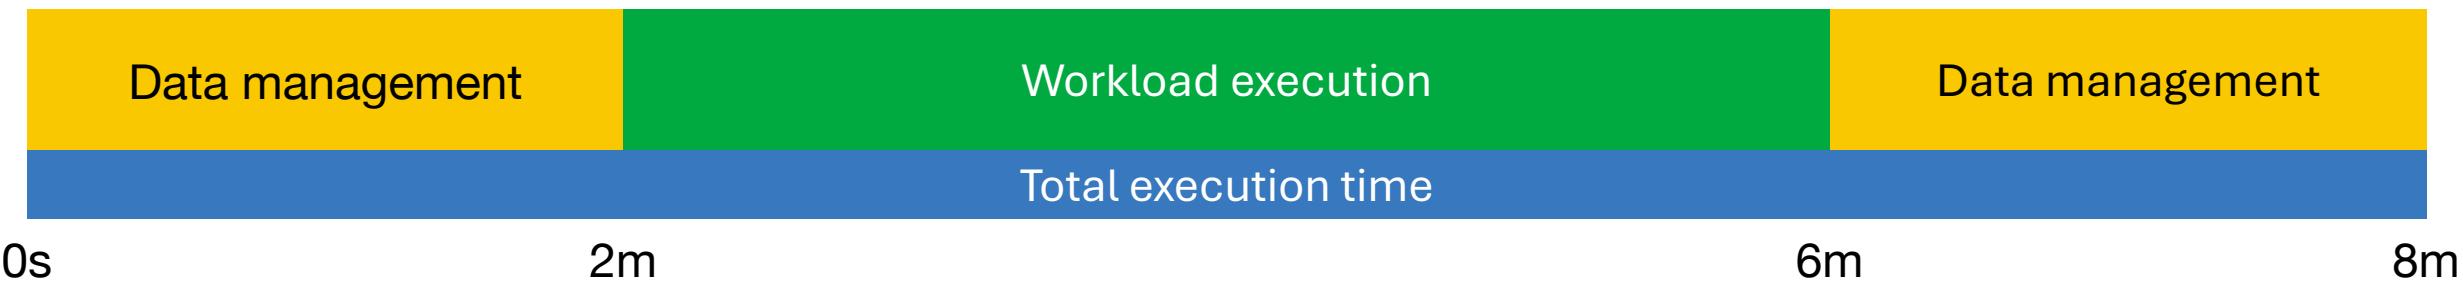
2 – Cluster Utilization – 3 4 5 6 7 8

Cluster Utilization

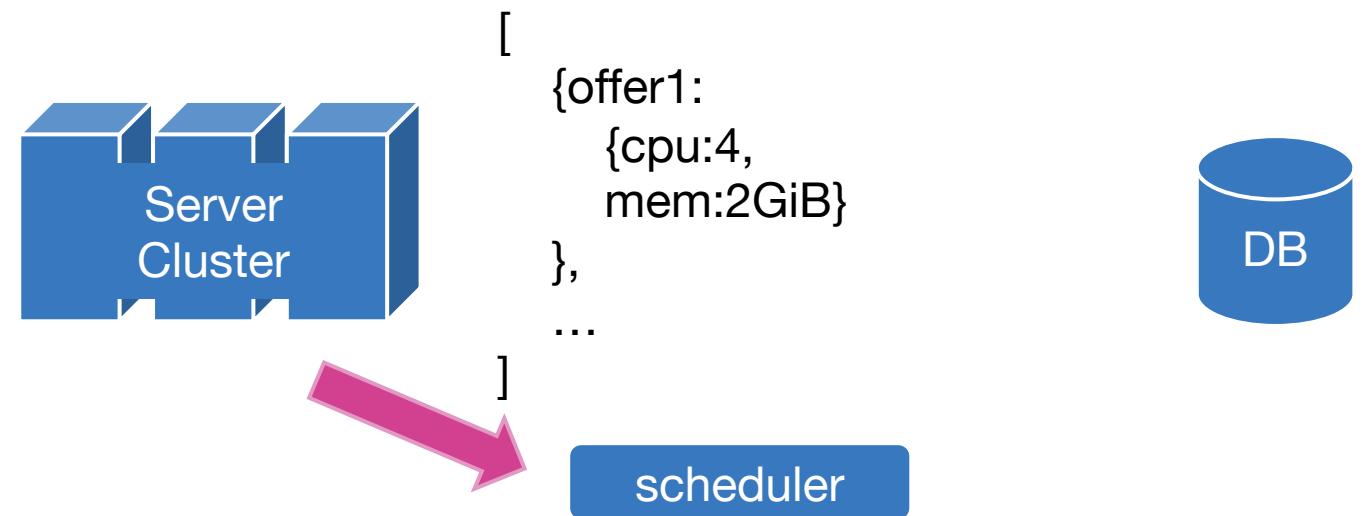


1 – Conflict

2 – Cluster Utilization – 3 4 5 6 7 8

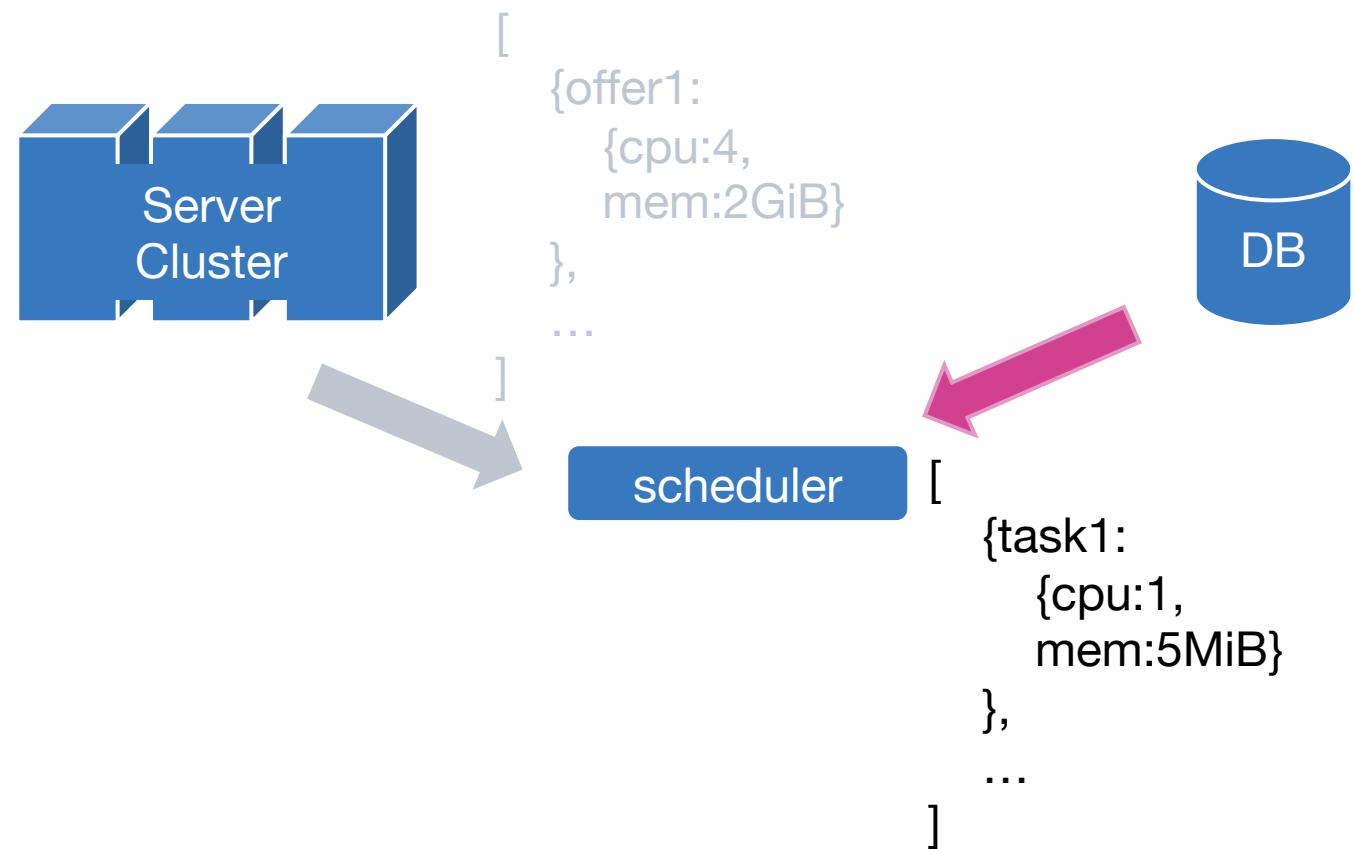


```
get_resource_offers_from_cluster
```

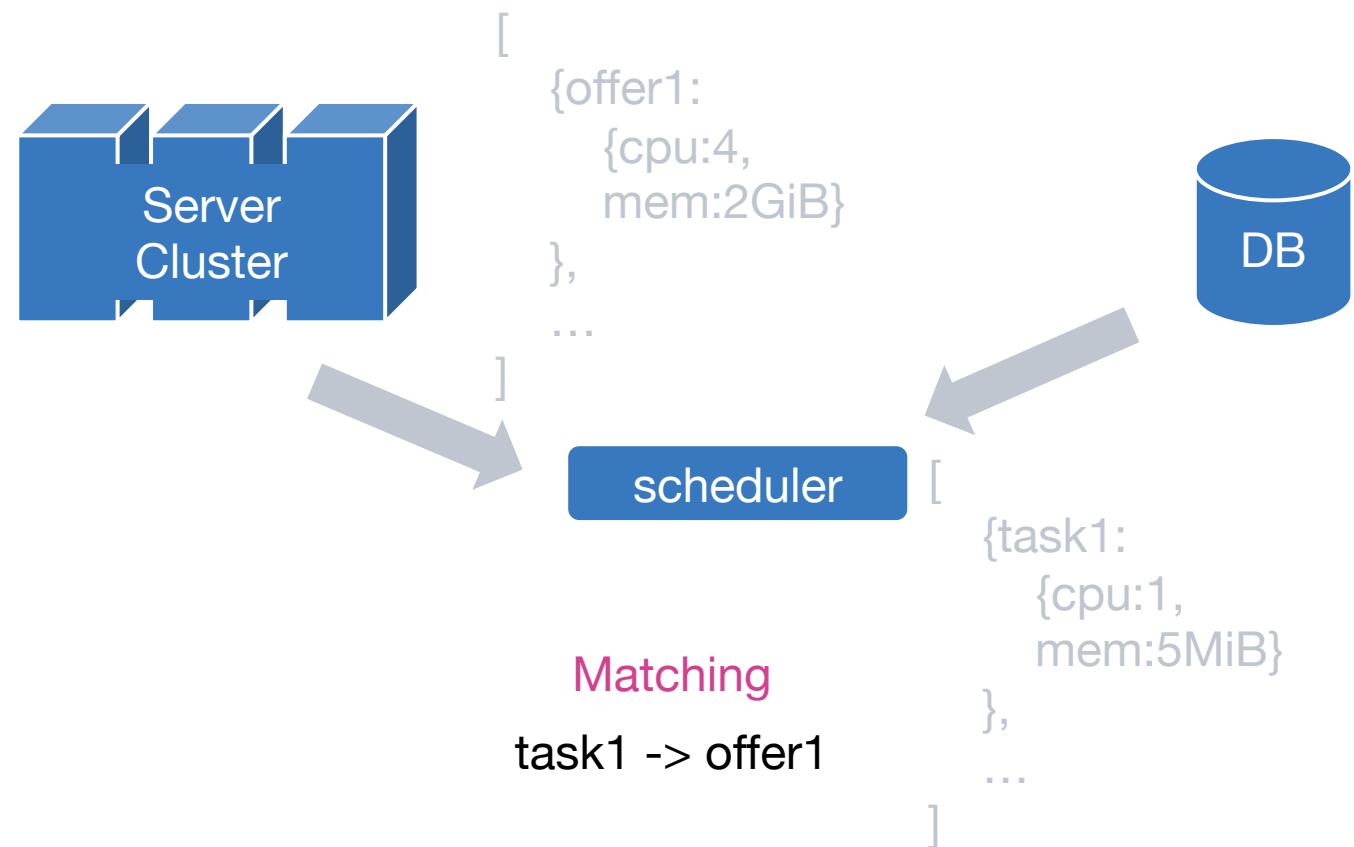


2 – Cluster Utilization

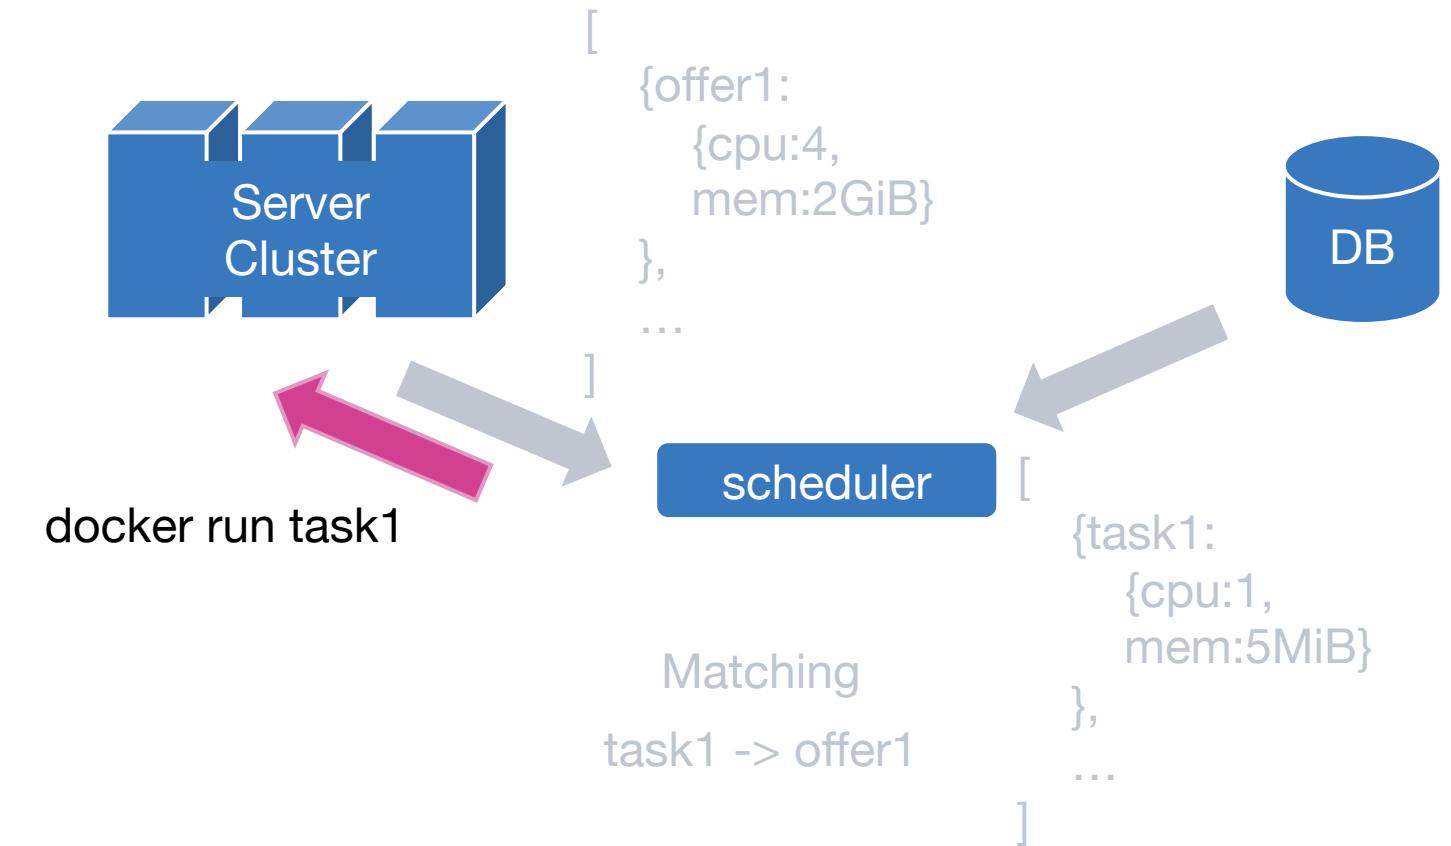
```
get_resource_offers_from_cluster  
getprocessable_tasks
```



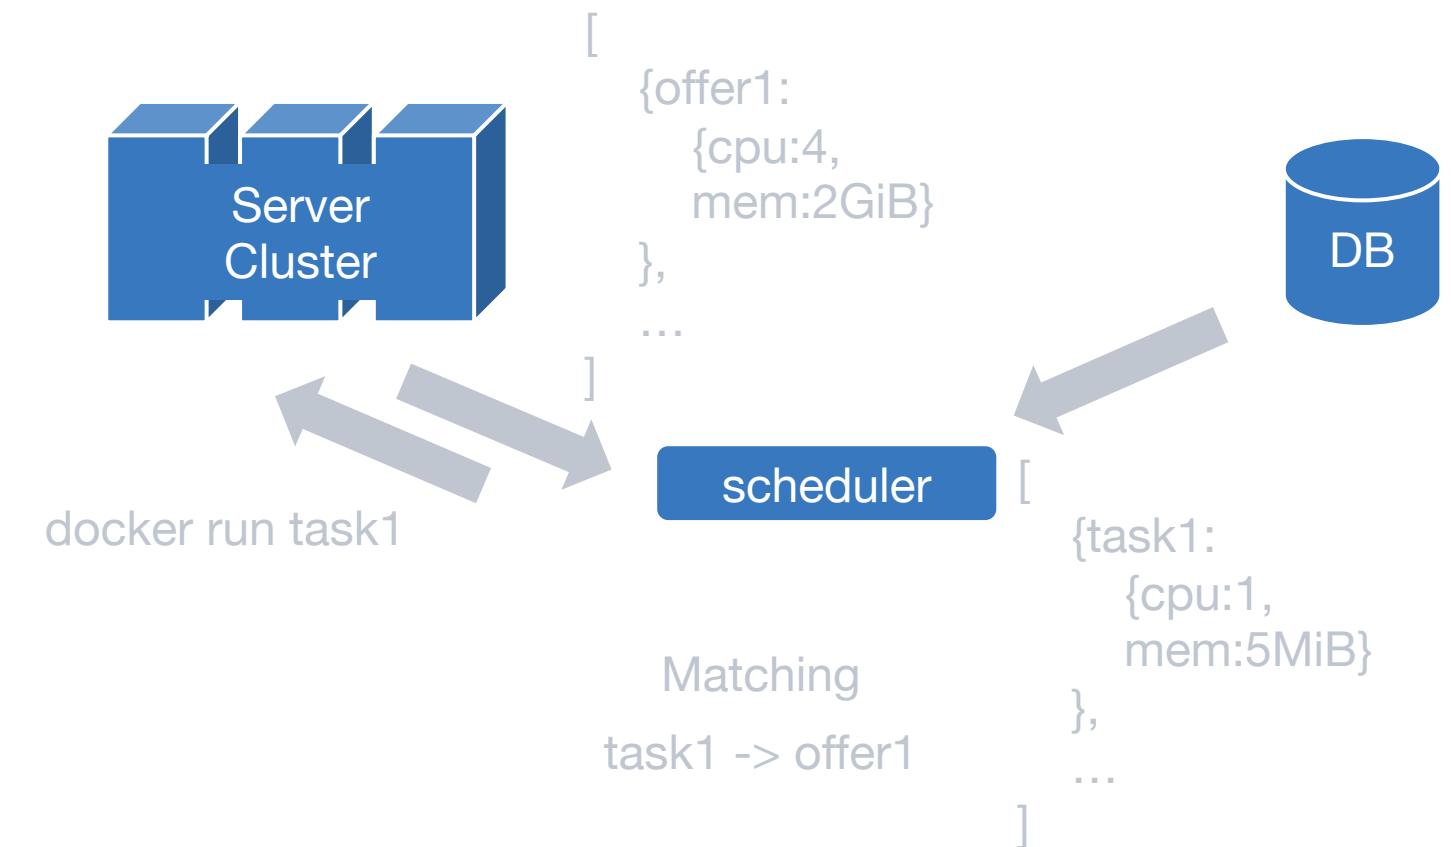
```
get_resource_offers_from_cluster  
getprocessable_tasks  
match_tasks_to_offers
```

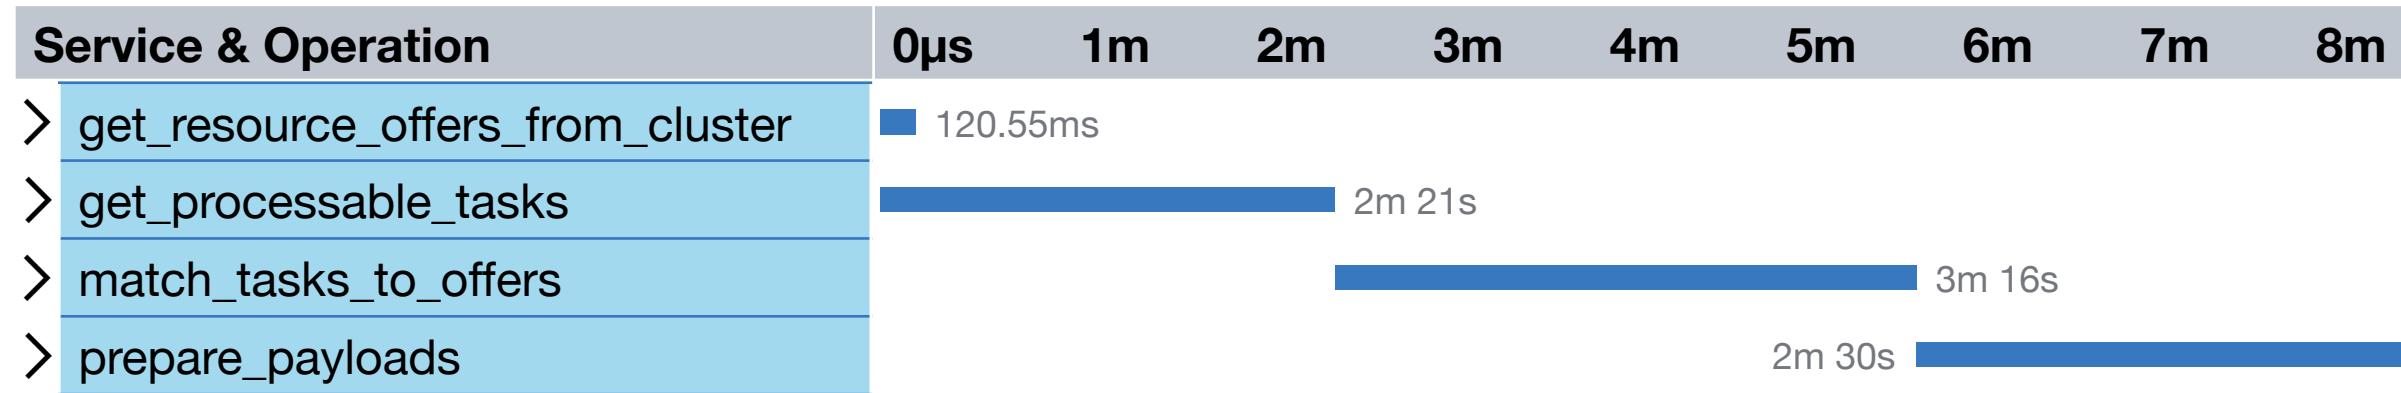


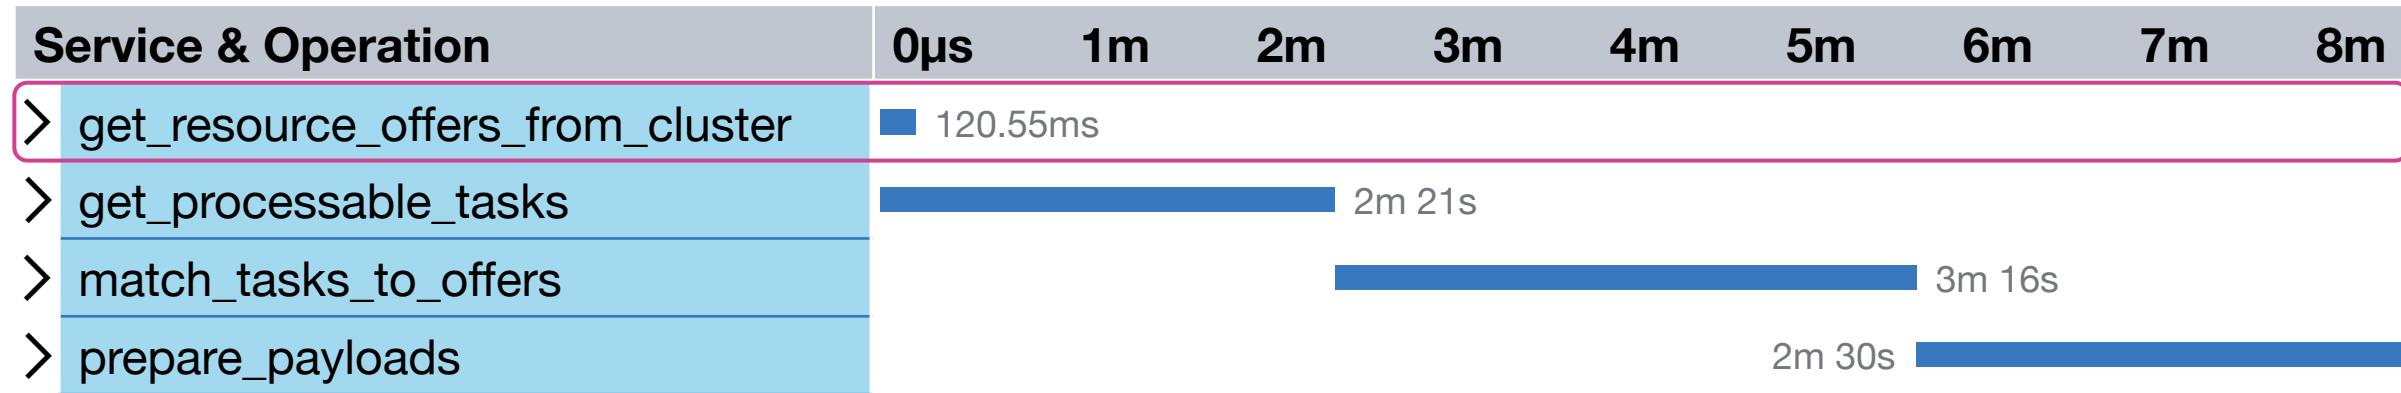
```
get_resource_offers_from_cluster  
getprocessable_tasks  
match_tasks_to_offers  
prepare_payloads
```

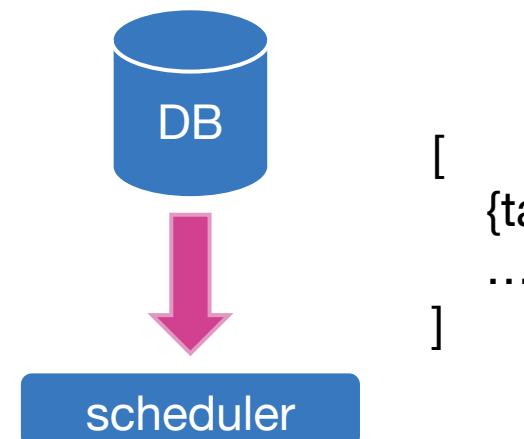
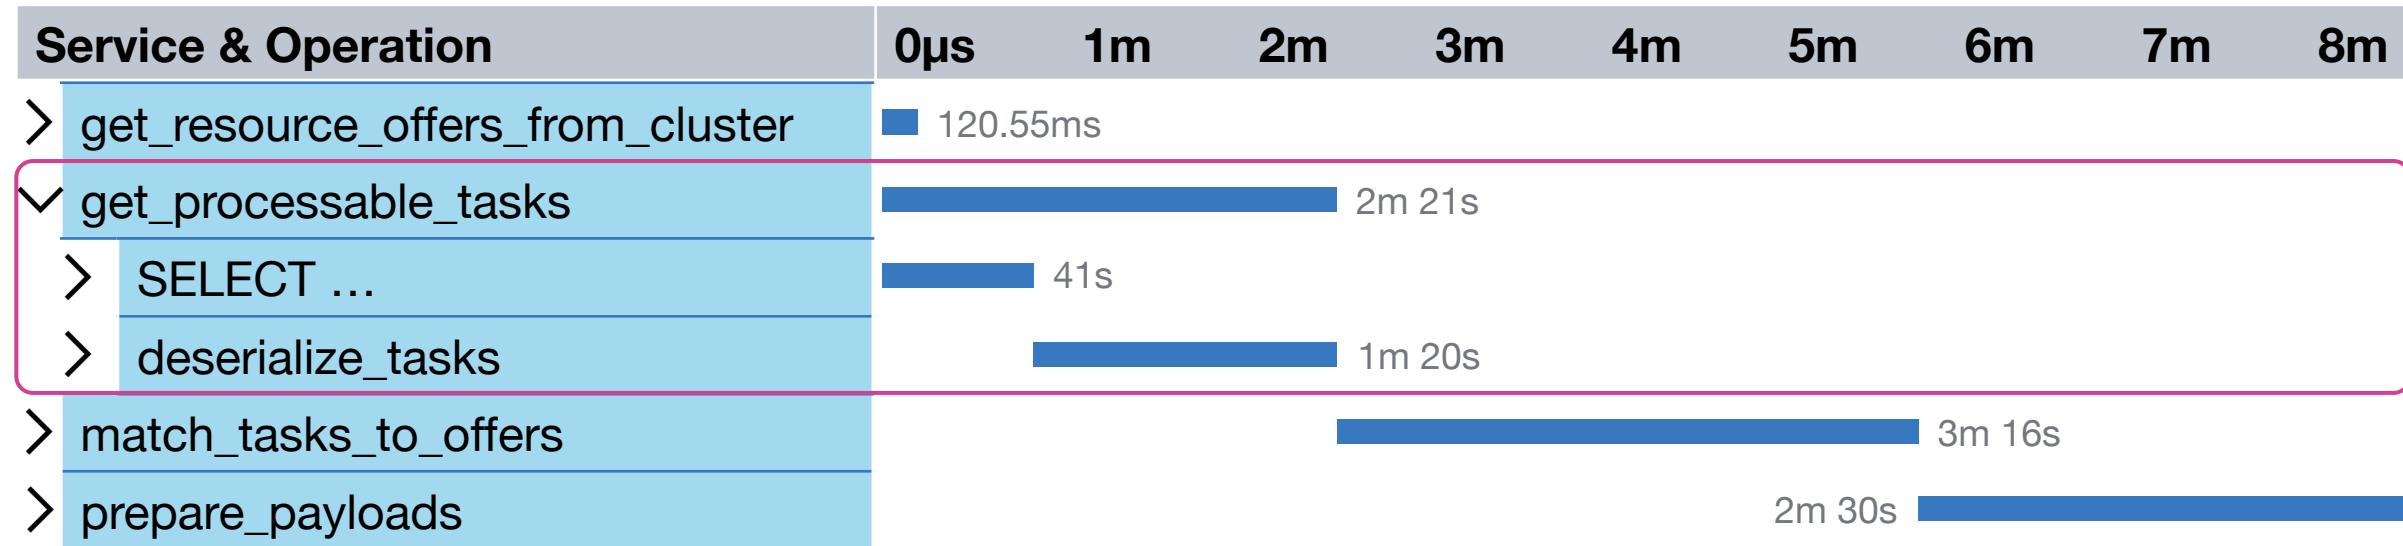


```
get_resource_offers_from_cluster  
getprocessable_tasks  
match_tasks_to_offers  
prepare_payloads
```

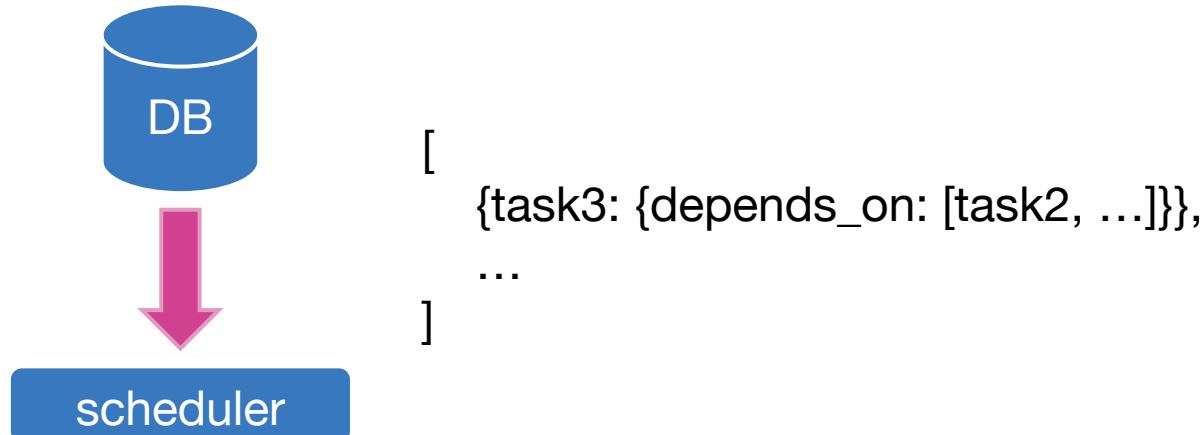
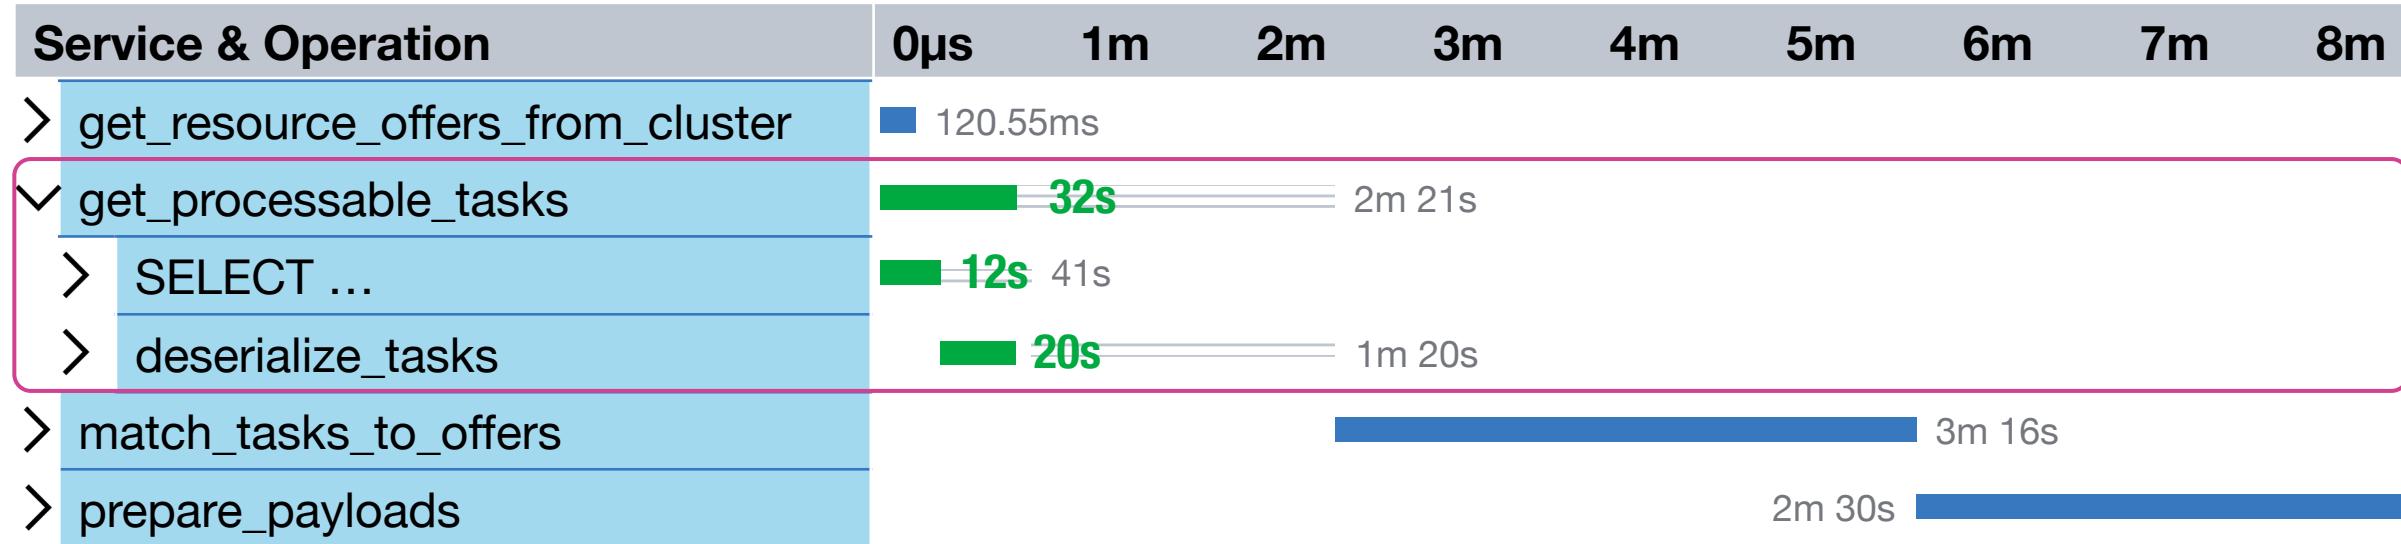


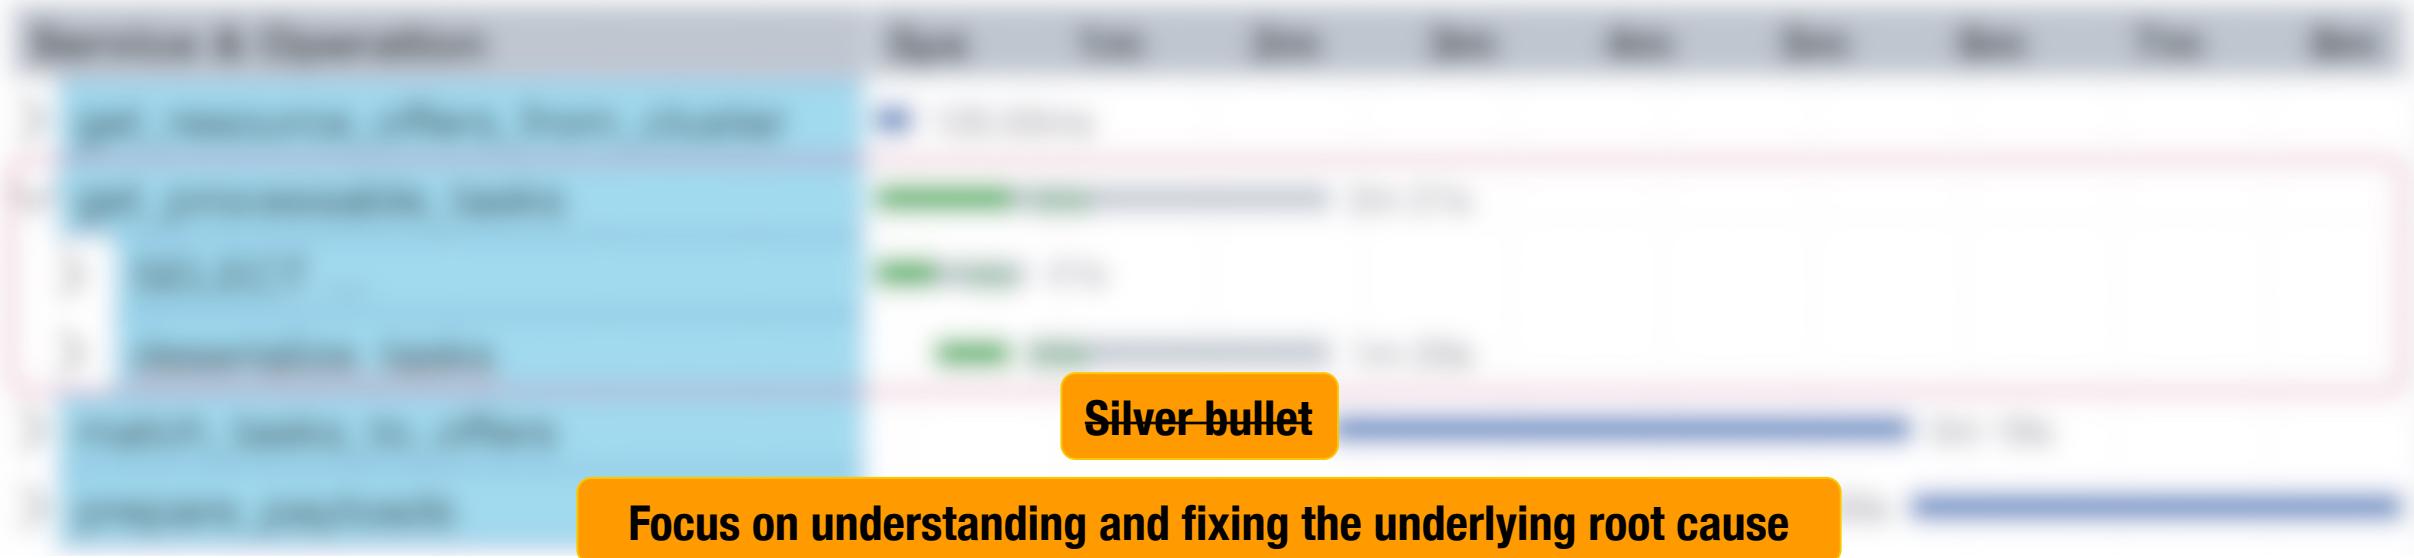






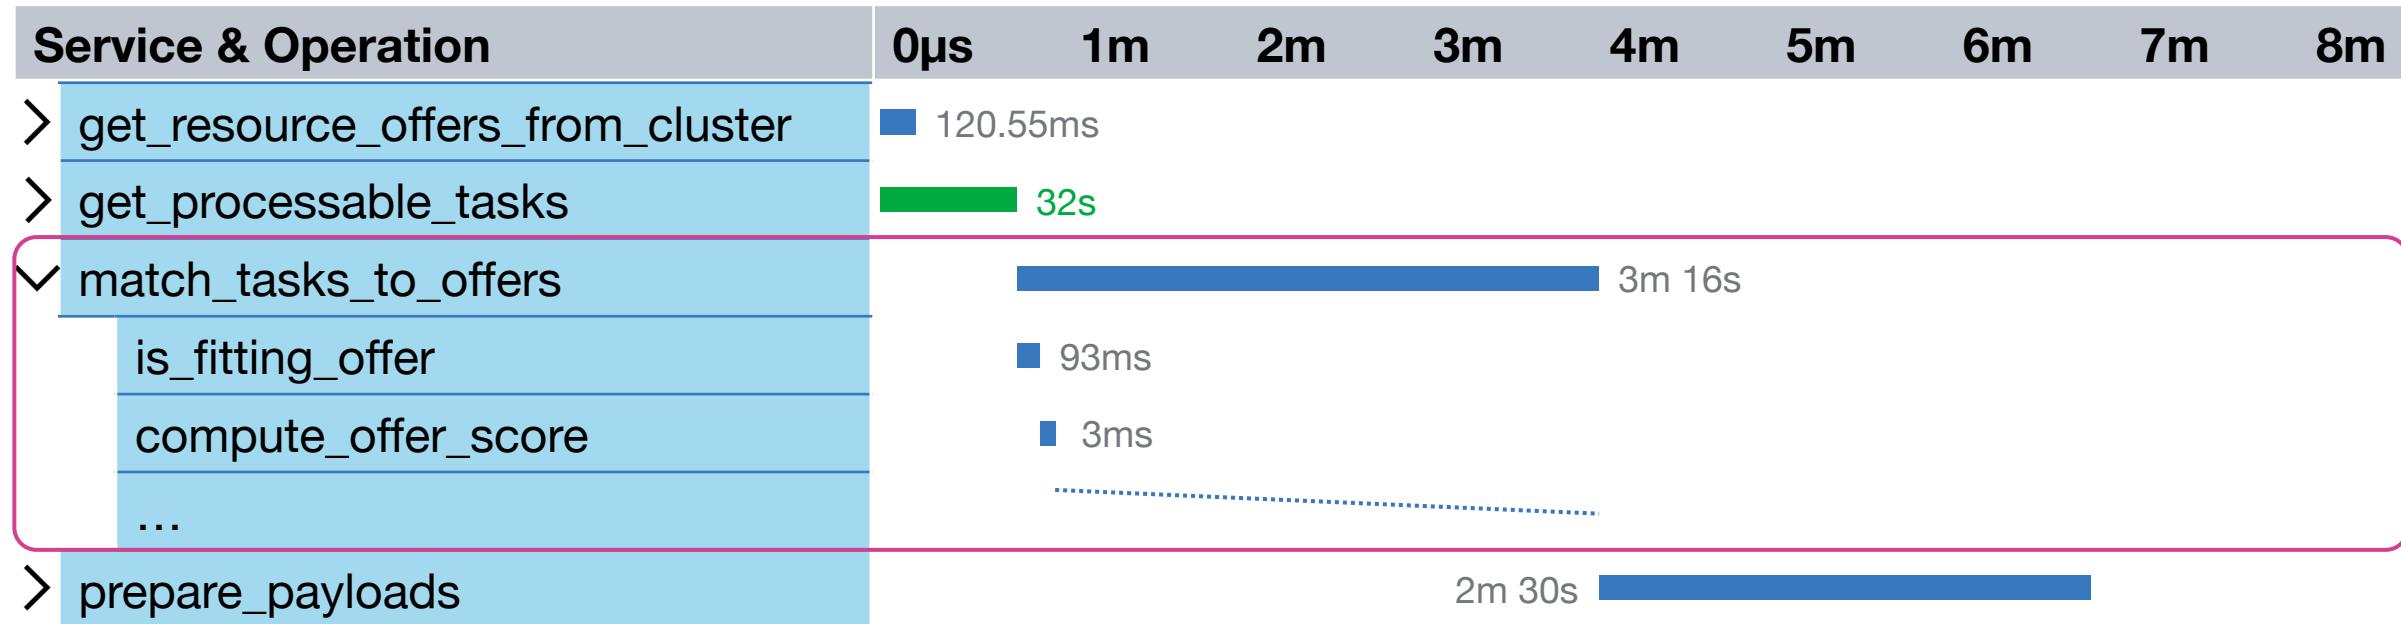
```
[  
  {task3: {depends_on: [{task2: ...}, ...]}},  
  ...  
(Note: The ellipsis '...' is shown between the first and second brackets.)]
```





Silver bullet

Focus on understanding and fixing the underlying root cause



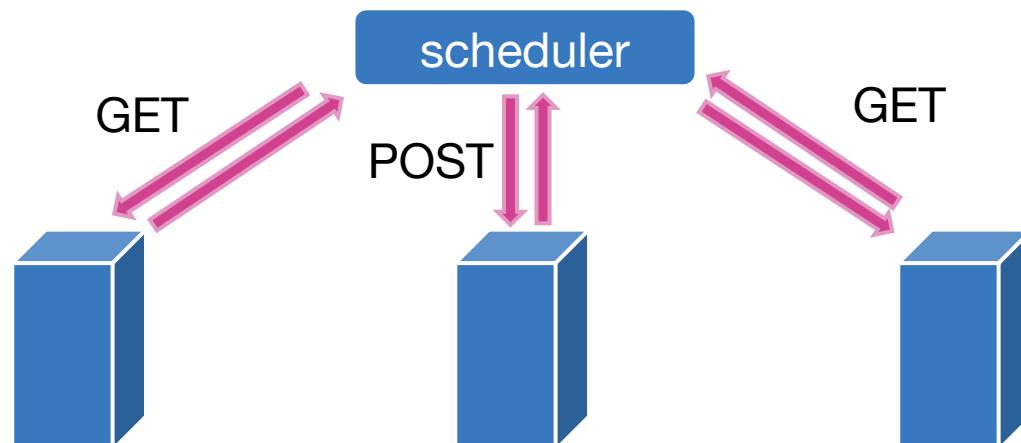
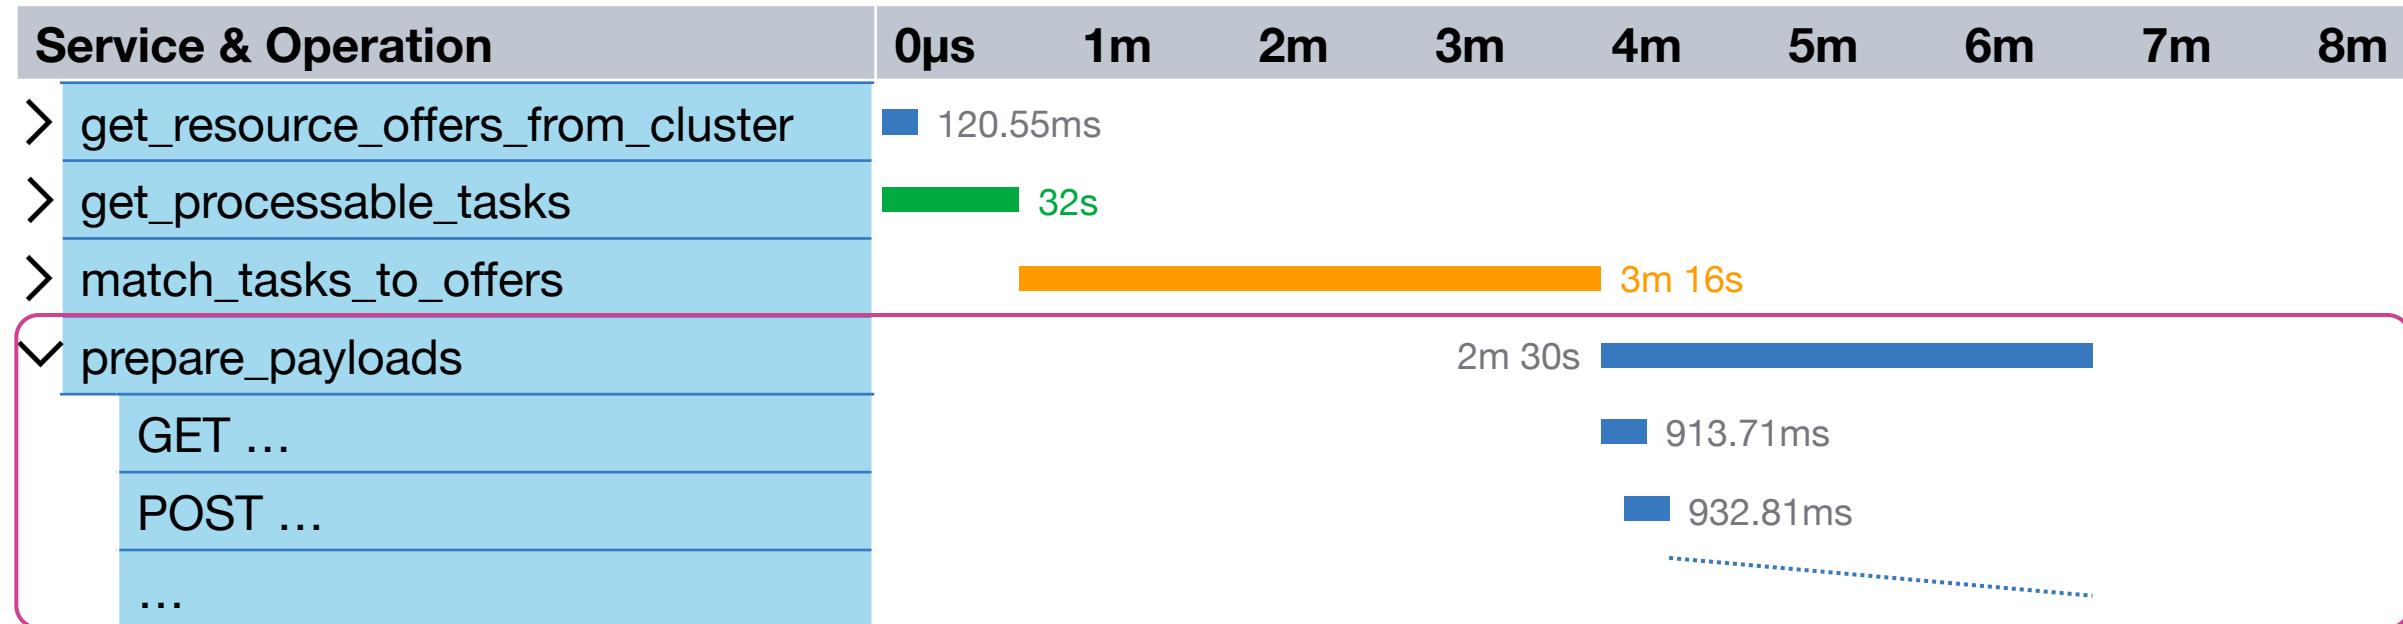
offer1: {cpu: 10, mem: 200GiB, ...}
 offer2: {cpu: 150, mem: 1.5TiB, ...}
 ...

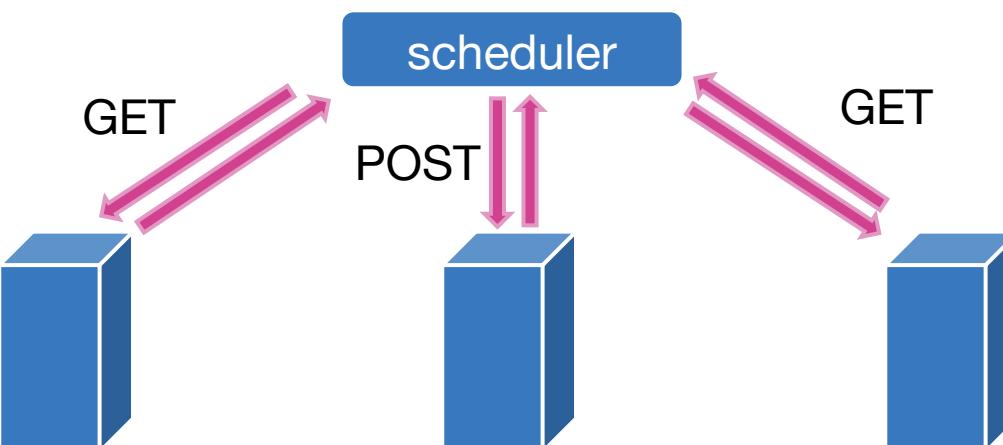
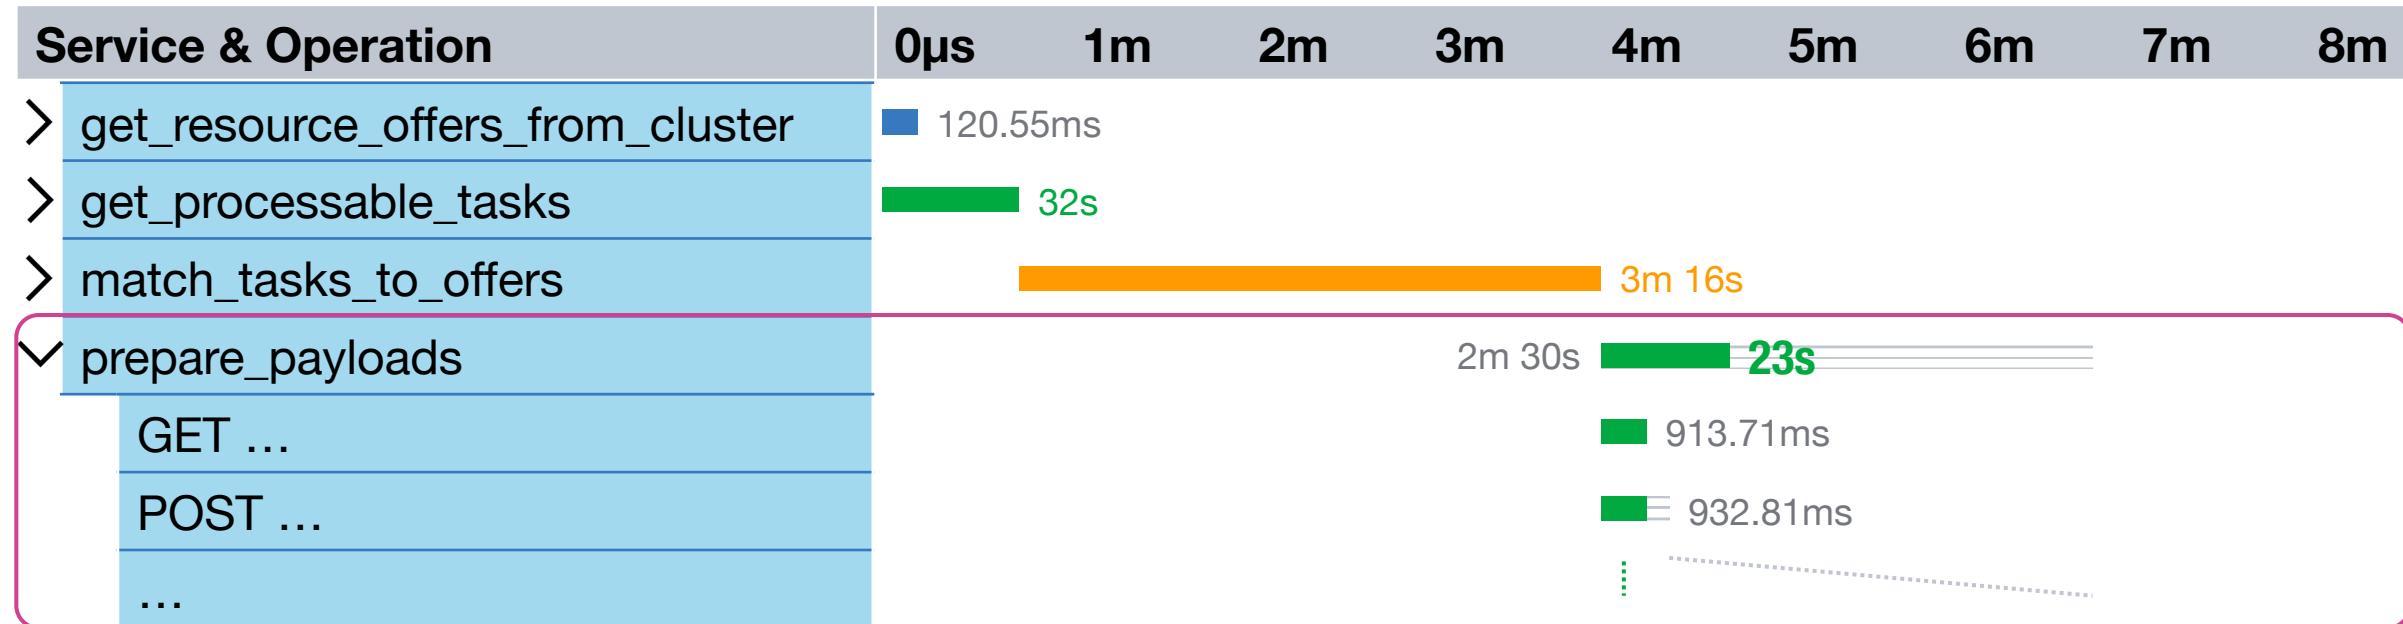


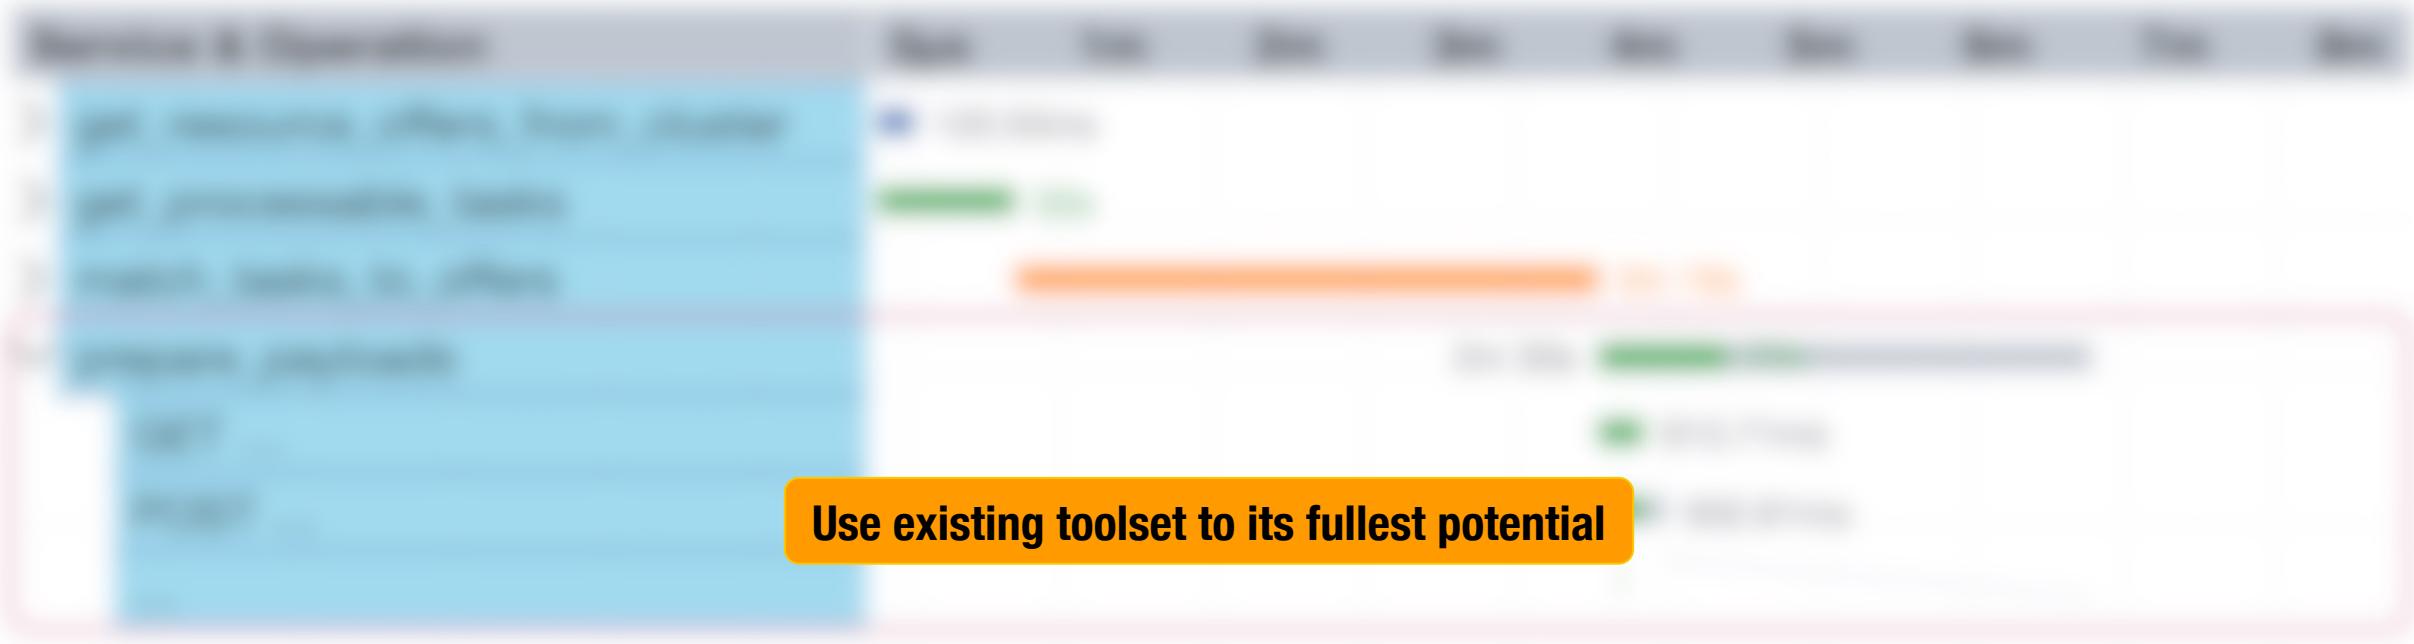
task1: {cpu: 5, mem: 100GiB, ...}
 task2: {cpu: 145, mem: 200GiB, ...}
 ...



offer1: []
 offer2: [task1, task2]







Use existing toolset to its fullest potential



```
11  def match_tasks_to_offers(  
12      tasks: list[Task], offers: list[Offer]  
13  ) -> dict[OfferName, TaskName]:  
14      """Returns a mapping from offer name to task name where the key is the name  
15      of the bestfit offer for the task.  
16      """  
17  
18      result = {}  
19  
20      for task in tasks:  
21          best_offer_name: tuple[OfferName, float] = None, None  
22  
23          for offer in sorted(offers):  
24              if is_fitting_offer(task, offer):  
25                  score = compute_score(task, offer)  
26                  if score < best_offer_name[1]:  
27                      best_offer_name = (offer, score)  
28  
29          if best_offer_name[0] is not None:  
30              result[best_offer_name[0].name] = task.name  
31  
32      return result
```



```
11  def match_tasks_to_offers(  
12      tasks: list[Task], offers: list[Offer]  
13  ) -> dict[OfferName, TaskName]:  
14      """Returns a mapping from offer name to task name where the key is the name  
15      of the bestfit offer for the task.  
16      """  
17      result = {}  
18      for task in tasks:  
19          best_index_score: tuple[OfferIndex | None, float] = (None, 31.0)  
20          for index, offer in enumerate(offers):  
21              if is_fitting_offer(task, offer):  
22                  score = compute_score(task, offer)  
23                  if score < best_index_score[1]:  
24                      best_index_score = (index, score)
```



```
11  def match_tasks_to_offers(  
12      tasks: list[Task], offers: list[Offer]  
13  ) -> dict[OfferName, TaskName]:  
14      """Returns a mapping from offer name to task name where the key is the name  
15      of the bestfit offer for the task.  
16      """  
17      result = {}  
18      for task in tasks:  
19          best_index_score: tuple[OfferIndex | None, float] = (None, 31.0)  
20          for index, offer in enumerate(offers):  
21              if is_fitting_offer(task, offer):  
22                  score = compute_score(task, offer)  
23                  if score < best_index_score[1]:  
24                      best_index_score = (index, score)  
25                  if (best_index := best_index_score[0]) is not None:  
26                      result[swap_remove(offers, best_index).name] = task.name  
27  
28      return result
```

| Implementation | Total time | is_fitting_offer time | compute_score time |
|------------------------|------------|-----------------------|--------------------|
| Production benchmark | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 83ms | 2ms |



```
128     def match_tasks_to_offers(  
129         pool_class: type[Pool], tasks: list[Task], offers: list[Offer]  
130     ) -> dict[OfferName, TaskName]:  
131         """Match the given tasks to the given offers"""  
132         pool = pool_class(os.cpu_count())  
  
133         result: dict[OfferName, TaskName] = {}  
134         for task in tasks:  
135             best_offer_score: tuple[OfferName | None, float] = (None, 0.0)  
136             for offer, score in offers:  
137                 pool.submit(get_offer_score, (task, offer))  
138             for offer in offers:  
139                 if score is not None and score > best_offer_score[1]:  
140                     best_offer_score = (offer, score)  
141             if best_offer_score[0] is None:  
142                 continue  
143             result[best_offer_score[0].name, best_offer_score[0].name] = task.name  
144         pool.shutdown()  
145         pool.join()  
146         return result
```



```
128     def match_tasks_to_offers(
129         pool_class: type[Pool], tasks: list[Task], offers: list[Offer]
130     ) -> dict[OfferName, TaskName]:
131         """Match the given tasks to the given offers"""
132         pool = pool_class(os.cpu_count())
133         result = {}
134         for task in tasks:
135             best_index_score: tuple[OfferIndex | None, float] = (None, 31.0)
136             for index, score in enumerate(
137                 pool.starmap(get_offer_score, [(task, offer) for offer in offers])
138             ):
139                 if score is not None and score < best_index_score[1]:
140                     best_index_score = (index, score)
141             if best_index_score[0] is None:
142                 continue
143             result[(swap_remove(offers, best_index_score[0]).name)] = task.name
144         pool.close()
145         pool.join()
146         return result
```

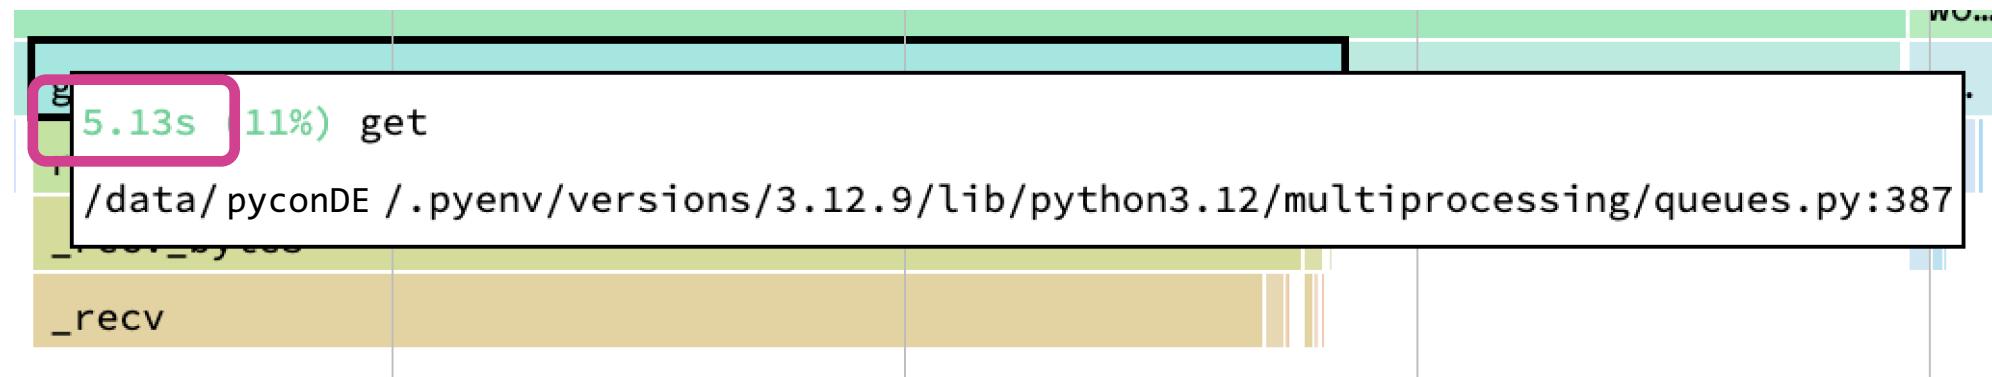
| Implementation | Total time | is_fitting_offer time | compute_score time |
|------------------------|------------|-----------------------|--------------------|
| Production benchmark | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 83ms | 2ms |
| Python multi threaded | 27s | 495ms | 29ms |

| Implementation | Total time | is_fitting_offer time | compute_score time |
|------------------------|------------|-----------------------|--------------------|
| Production benchmark | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 83ms | 2ms |
| Python multi threaded | 27s | 495ms | 29ms |



| Implementation | Total time | is_fitting_offer time | compute_score time |
|-------------------------|------------|-----------------------|--------------------|
| Production benchmark | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 83ms | 2ms |
| Python multi threaded | 27s | 495ms | 29ms |
| Python multi processing | 70s | 50ms | 2ms |

| Implementation | Total time | is_fitting_offer time | compute_score time |
|-------------------------|------------|-----------------------|--------------------|
| Production benchmark | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 83ms | 2ms |
| Python multi threaded | 27s | 495ms | 29ms |
| Python multi processing | 70s | 50ms | 2ms |

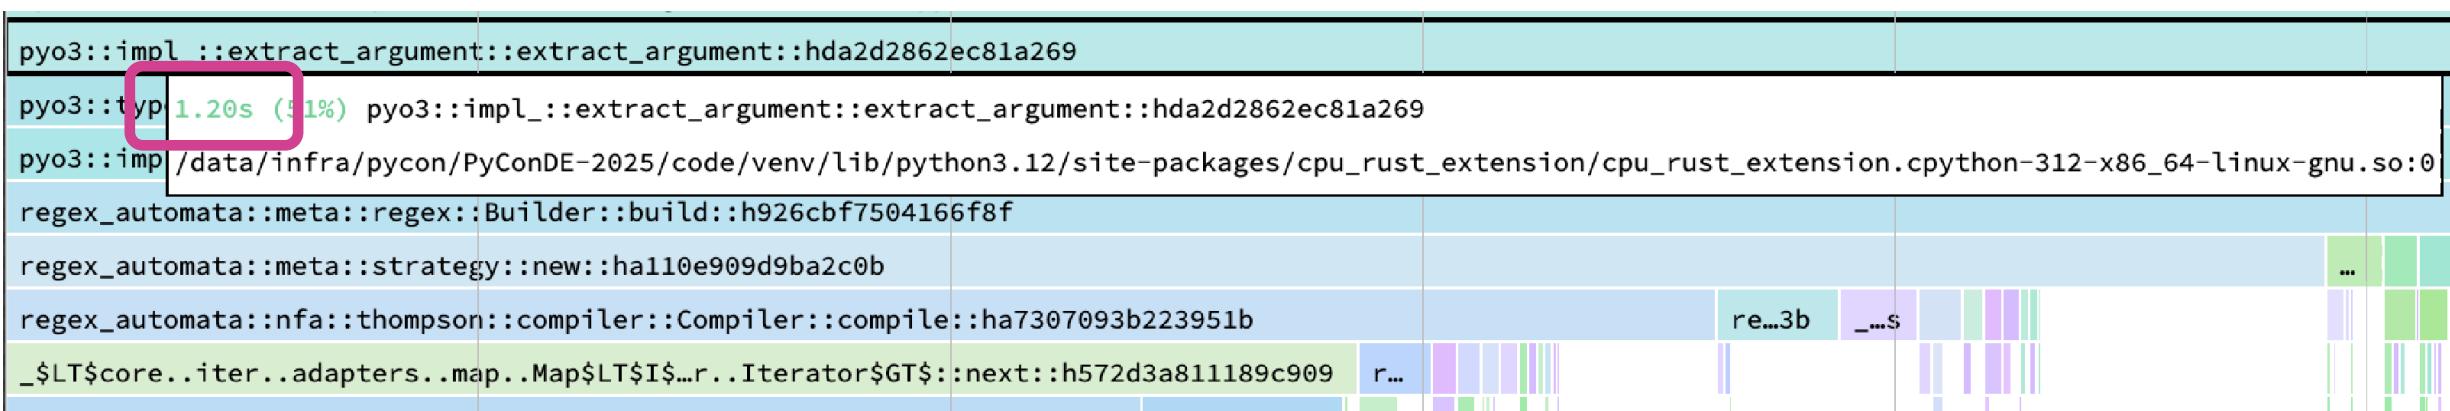


```
11 fn match_tasks_to_offers(tasks: Vec<Task>, mut offers: Vec<Offer>) -> HashMap<String, String> {
12     let mut result = HashMap::with_capacity(offers.len());
13     let mut best_index_score: (Option<usize>, f64);
14     for task in tasks.iter() {
15         best_index_score = (None, 31.0);
16         for (index, offer) in offers.iter().enumerate() {
17             if is_fitting_offer(task, offer) {
18                 let score = compute_score(task, offer);
19                 if score < best_index_score.1 {
20                     best_index_score = (Some(index), score);
21                 }
22             }
23         }
24         if let Some(best_index) = best_index_score.0 {
25             result.insert(offers.swap_remove(best_index).name, task.name.clone());
26         }
27     }
28     result
29 }
```



| Implementation | Total time | is_fitting_offer time | compute_score time |
|-------------------------|------------|-----------------------|--------------------|
| Production benchmark | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 83ms | 2ms |
| Python multi threaded | 27s | 495ms | 29ms |
| Python multi processing | 70s | 50ms | 2ms |
| Rust single threaded | 2s | 20ms | 150µs |

| Implementation | Total time | is_fitting_offer time | compute_score time |
|-------------------------|------------|-----------------------|--------------------|
| Production benchmark | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 83ms | 2ms |
| Python multi threaded | 27s | 495ms | 29ms |
| Python multi processing | 70s | 50ms | 2ms |
| Rust single threaded | 2s | 20ms | 150µs |



| Implementation | Total time | is_fitting_offer time | compute_score time |
|-------------------------|------------|-----------------------|--------------------|
| Production benchmark | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 83ms | 2ms |
| Python multi threaded | 27s | 495ms | 29ms |
| Python multi processing | 70s | 50ms | 2ms |
| Rust single threaded | 2s | 20ms | 150µs |
| Rust multi threaded | 3s | 20ms | 150µs |

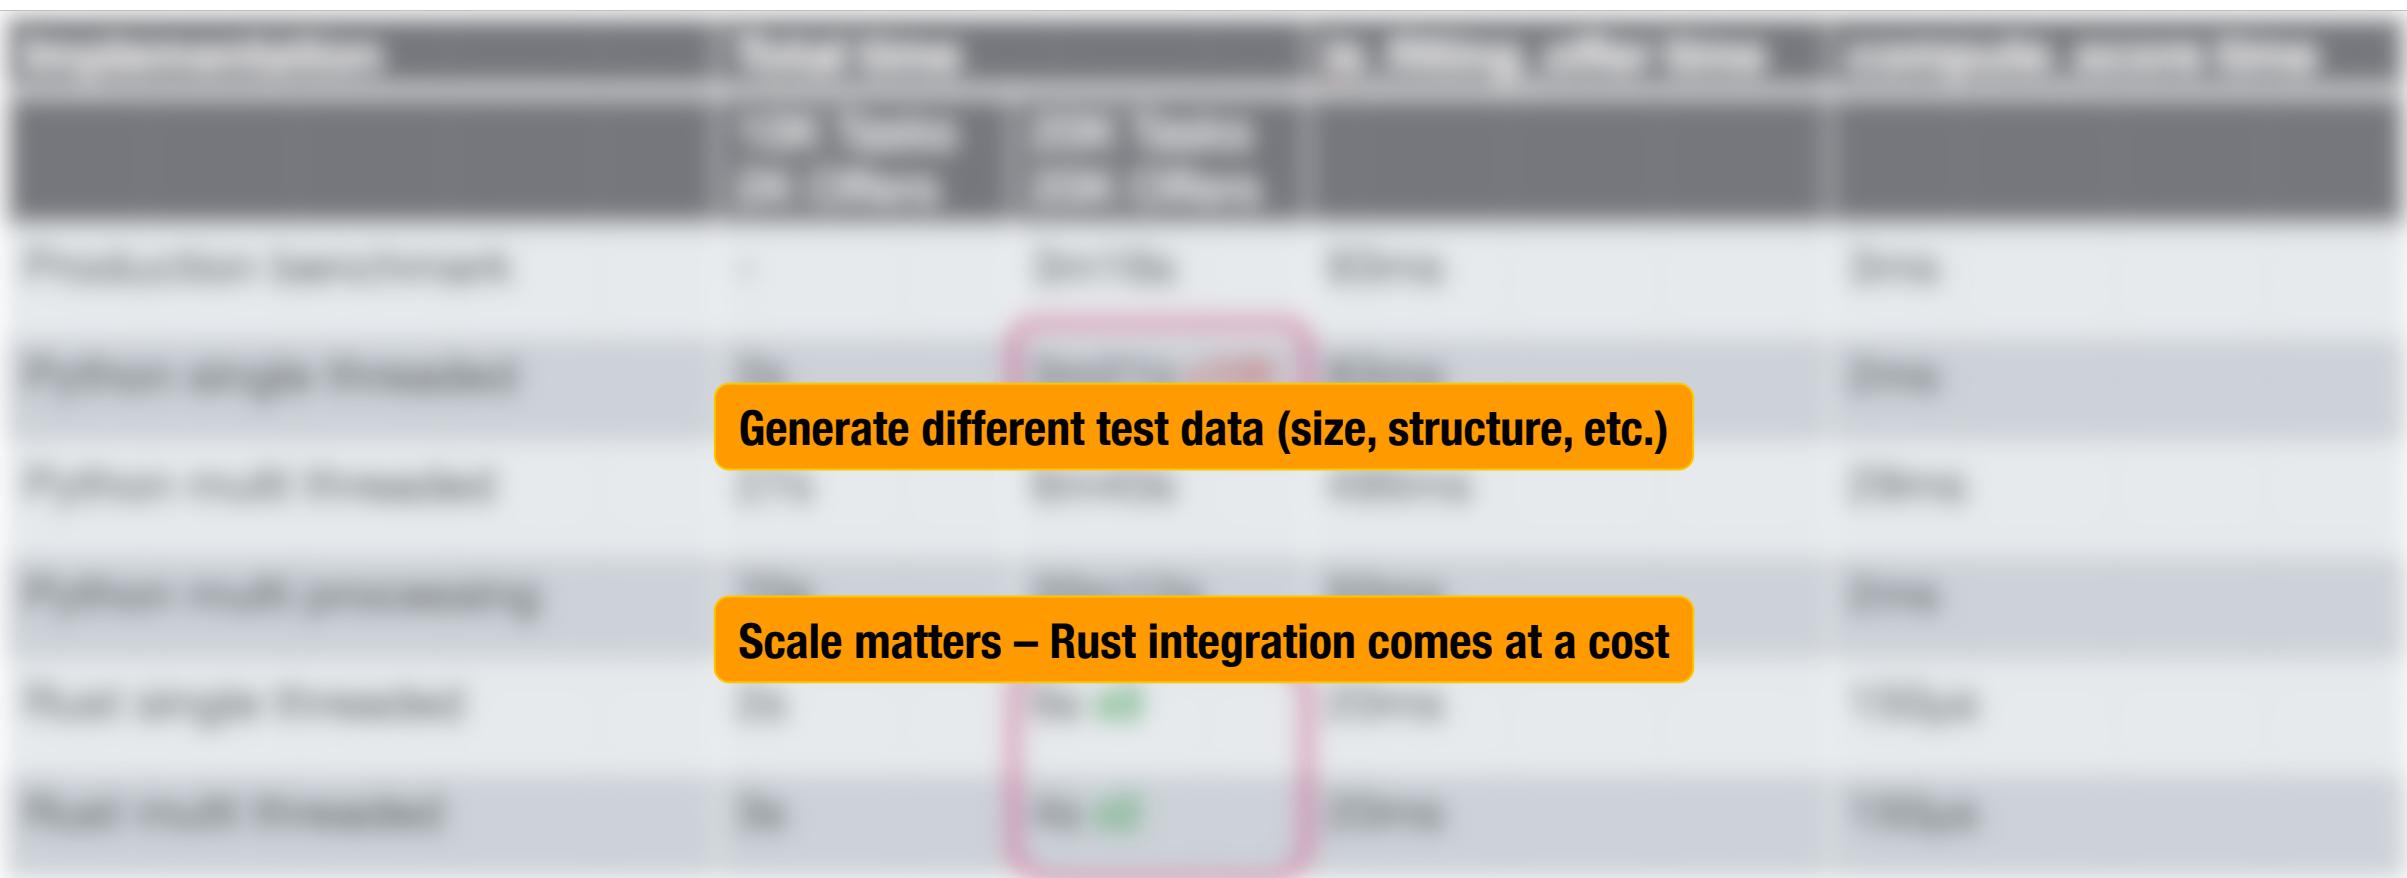
| Implementation | Total time | | is_fitting_offer time | compute_score time |
|-------------------------|------------------------|-------------------------|-----------------------|--------------------|
| | 10K Tasks 2K Offers | 20K Tasks 20K Offers | | |
| Production benchmark | - | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | | 83ms | 2ms |
| Python multi threaded | 27s | | 495ms | 29ms |
| Python multi processing | 70s | | 50ms | 2ms |
| Rust single threaded | 2s | | 20ms | 150µs |
| Rust multi threaded | 3s | | 20ms | 150µs |

| Implementation | Total time | | is_fitting_offer time | compute_score time |
|-------------------------|------------------------|-------------------------|-----------------------|--------------------|
| | 10K Tasks 2K Offers | 20K Tasks 20K Offers | | |
| Production benchmark | - | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 3m21s | 83ms | 2ms |
| Python multi threaded | 27s | 6m40s | 495ms | 29ms |
| Python multi processing | 70s | 20m12s | 50ms | 2ms |
| Rust single threaded | 2s | 6s | 20ms | 150µs |
| Rust multi threaded | 3s | 4s | 20ms | 150µs |

| Implementation | Total time | | is_fitting_offer time | compute_score time |
|-------------------------|------------------------|-------------------------|-----------------------|--------------------|
| | 10K Tasks 2K Offers | 20K Tasks 20K Offers | | |
| Production benchmark | - | 3m16s | 93ms | 3ms |
| Python single threaded | 2s | 3m21s | 83ms | 2ms |
| Python multi threaded | 27s | 6m40s | 495ms | 29ms |
| Python multi processing | 70s | 20m12s | 50ms | 2ms |
| Rust single threaded | 2s | 6s | 20ms | 150µs |
| Rust multi threaded | 3s | 4s | 20ms | 150µs |

The diagram illustrates the performance gains of the Rust implementations compared to their Python counterparts. A green circle highlights the Rust single-threaded implementation's speedup of 30x over Python single-threaded. A green arrow highlights the Rust multi-threaded implementation's speedup of 49x over Python single-threaded.

| Implementation | Total time | | is_fitting_offer time | compute_score time |
|-------------------------|---|-------------------------|-----------------------|--------------------|
| | 10K Tasks 2K Offers | 20K Tasks 20K Offers | | |
| Production benchmark | - | 3m16s | 93ms | 3ms |
| Python single threaded | 2s  3m21s | 83ms | 2ms | |
| Python multi threaded | 27s | 6m40s | 495ms | 29ms |
| Python multi processing | 70s | 20m12s | 50ms | 2ms |
| Rust single threaded | 2s  6s | 20ms | 150µs | |
| Rust multi threaded | 3s  4s | 20ms | 150µs | |



Generate different test data (size, structure, etc.)

Scale matters – Rust integration comes at a cost

| Implementation | Total time (+increase) | Background thread operations/ms |
|-------------------------|------------------------|---------------------------------|
| Python single threaded | 3m51s (+ 30s) | 61 |
| Python multi threaded | 7m40s (+ 1m) | 51 |
| Python multi processing | 20m12s (+ 0s) | 79 |
| Rust single threaded | 6s (+ 0s) | 0 |
| Rust multi threaded | 4s (+ 0s) | 0 |

| Implementation | Total time (+increase) | Background thread operations/ms |
|-------------------------|------------------------|---------------------------------|
| Python single threaded | 3m51s (+ 30s) | 61 |
| Python multi threaded | 7m40s (+ 1m) | 51 |
| Python multi processing | 20m12s (+ 0s) | 79 |
| Rust single threaded | 6s (+ 0s) | 0 |
| Rust multi threaded | 4s (+ 0s) | 0 |

| Implementation | Total time (+increase) | Background thread operations/ms |
|-------------------------|------------------------|---------------------------------|
| Python single threaded | 3m51s (+ 30s) | 61 |
| Python multi threaded | 7m40s (+ 1m) | 51 |
| Python multi processing | 20m12s (+ 0s) | 79 |
| Rust single threaded | 6s (+ 0s) | 0 |
| Rust multi threaded | 4s (+ 0s) | 0 |

```

32  32 #[pyfunction]
33  33 #[pyo3(name = "match_single_threaded")]
34  pub fn main(tasks: Vec<Task>, offers: Vec<Offer>) -> HashMap<String, String> {
35      match_tasks_to_offers(tasks, offers)
34  pub fn main(py: Python, tasks: Vec<Task>, offers: Vec<Offer>) -> HashMap<String, String> {
35      py.allow_threads(move || match_tasks_to_offers(tasks, offers))
36  }

```



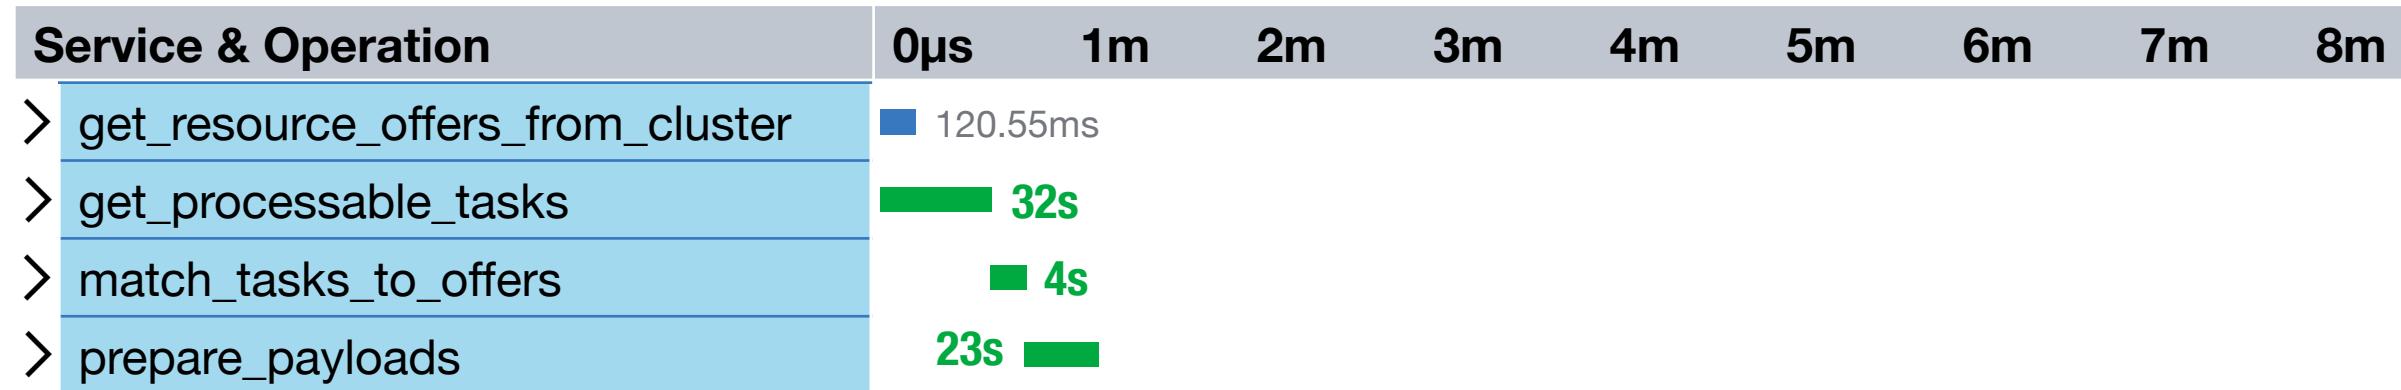
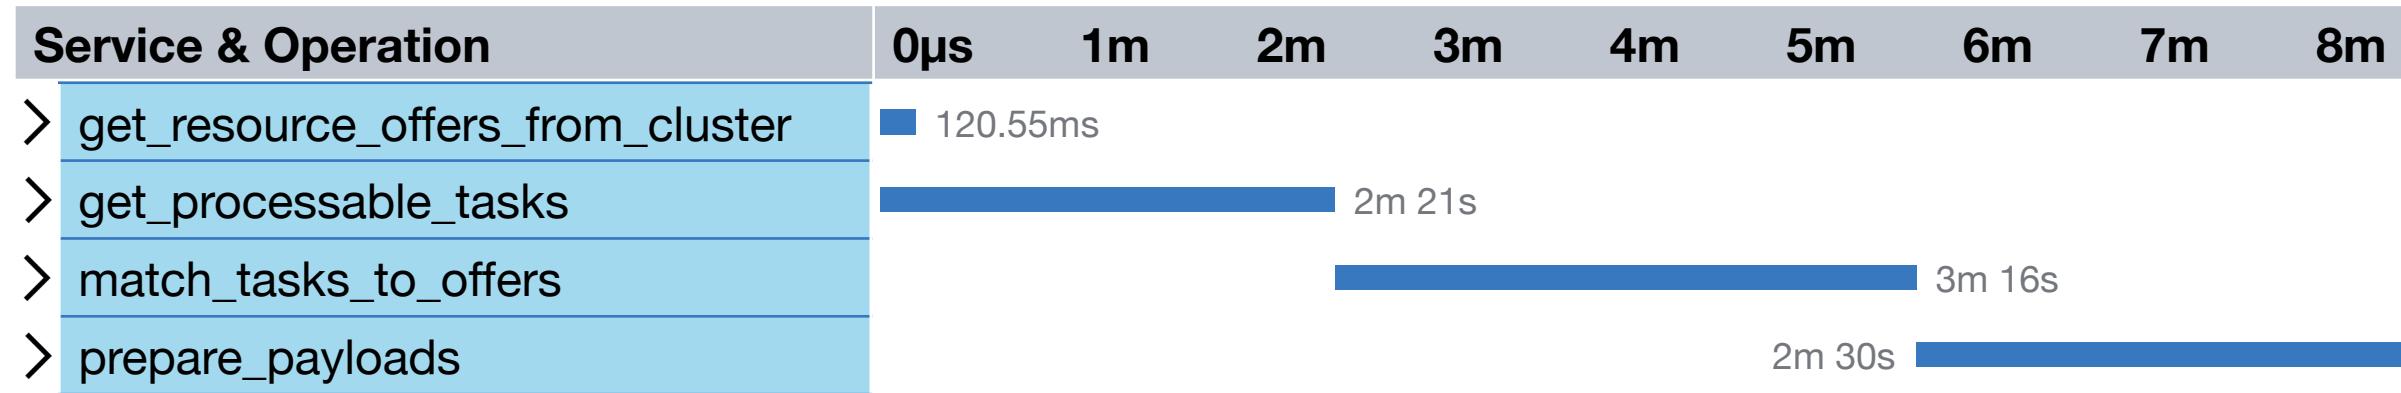
| Implementation | Total time (+increase) | Background thread operations/ms |
|-------------------------|------------------------|---------------------------------|
| Python single threaded | 3m51s (+ 30s) | 61 |
| Python multi threaded | 7m40s (+ 1m) | 51 |
| Python multi processing | 20m12s (+ 0s) | 79 |
| Rust single threaded | 6s (+ 50ms) | 80 |
| Rust multi threaded | 4s (+ 50ms) | 79 |

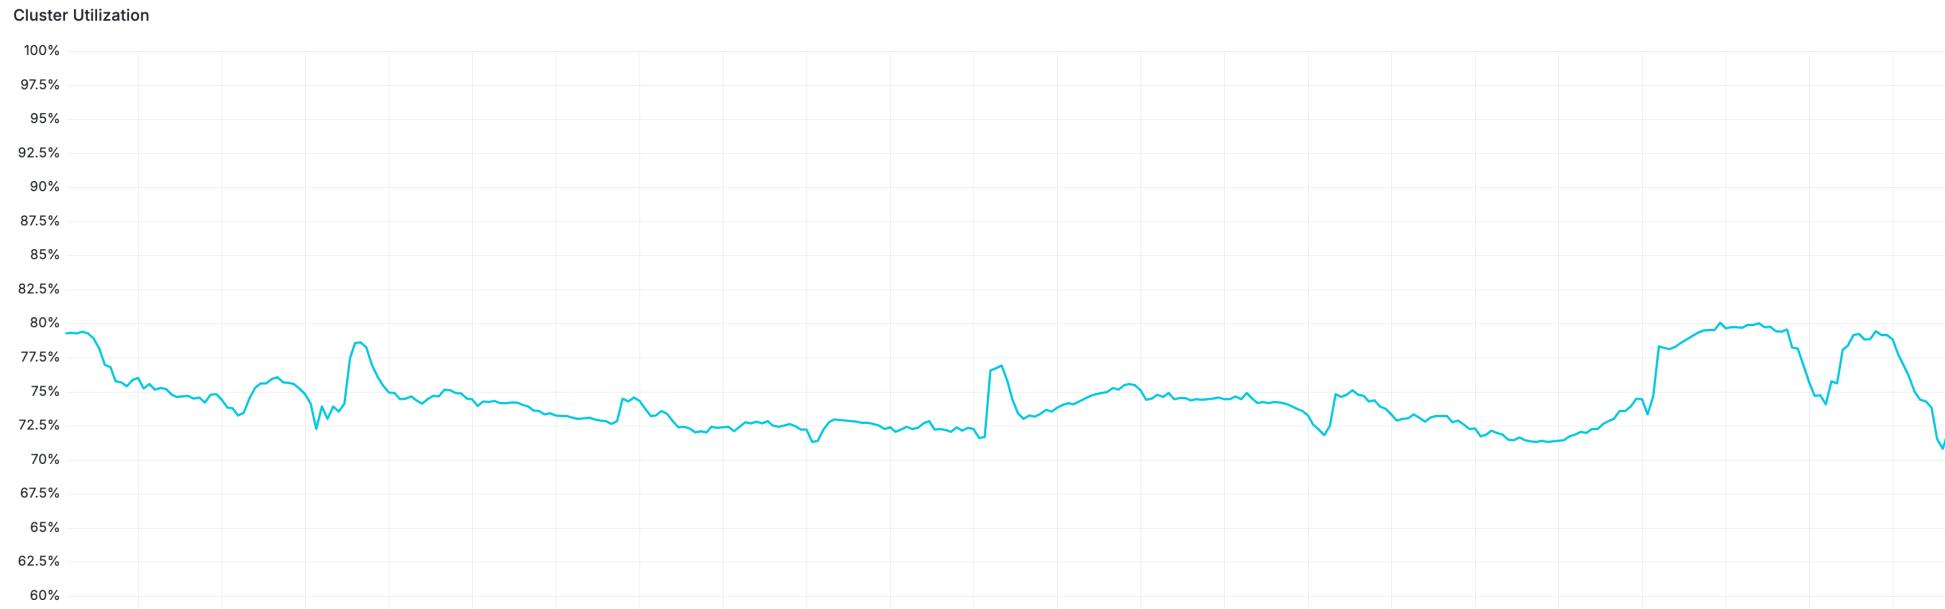
```

32  32 #[pyfunction]
33  33 #[pyo3(name = "match_single_threaded")]
34  pub fn main(tasks: Vec<Task>, offers: Vec<Offer>) -> HashMap<String, String> {
35      match_tasks_to_offers(tasks, offers)
34  pub fn main(py: Python, tasks: Vec<Task>, offers: Vec<Offer>) -> HashMap<String, String> {
35      py.allow_threads(move || match_tasks_to_offers(tasks, offers))
36  }

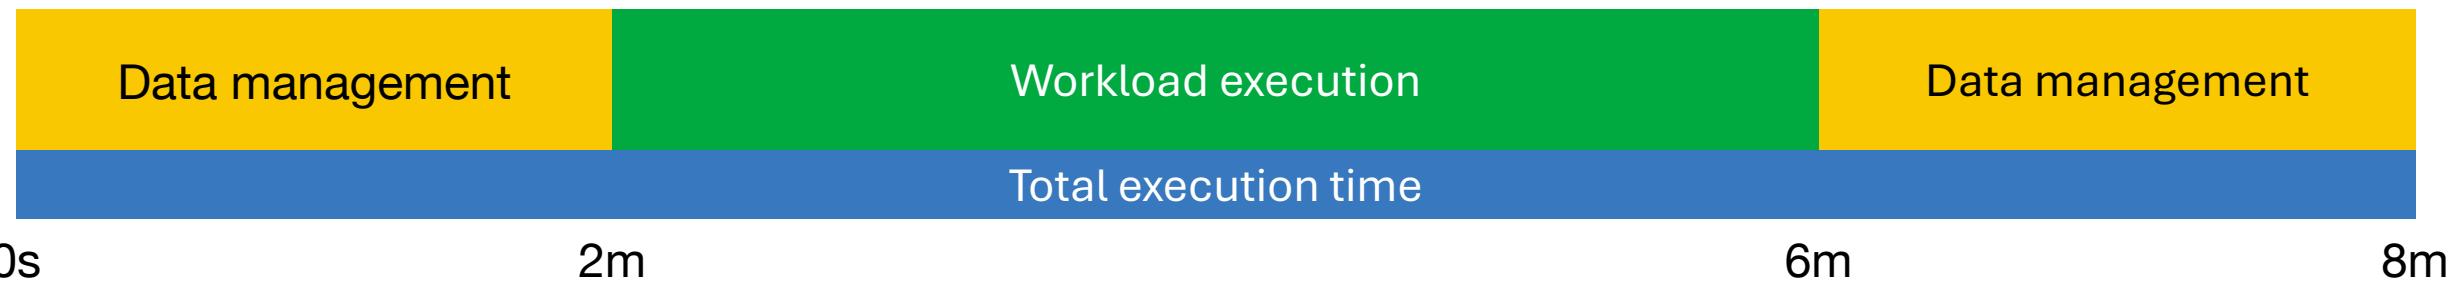
```







| Service & Operation | 0µs | 1m | 2m | 3m | 4m | 5m | 6m | 7m | 8m |
|------------------------------------|----------|-----|----|----|----|----|----|----|----|
| > get_resource_offers_from_cluster | 120.55ms | | | | | | | | |
| > getprocessable_tasks | | 32s | | | | | | | |
| > match_tasks_to_offers | | | 4s | | | | | | |
| > prepare_payloads | 23s | | | | | | | | |



| "External factor" | Used | Available |
|--------------------------|----------|-----------|
| Local disk throughput | 482 MB/s | 1.2 GB/s |
| Cloud storage throughput | 5 GB/s | >100 GB/s |
| Network bandwidth | 5 GB/s | 10 GB/s |

| "External factor" | Used | Available |
|--------------------------|----------|-----------|
| Local disk throughput | 482 MB/s | 1.2 GB/s |
| Cloud storage throughput | 5 GB/s | >100 GB/s |
| Network bandwidth | 5 GB/s | 10 GB/s |

```
6  ✓ def main(source: str, target: str) -> None:  
7      """Copy source to target path"""  
8  ✓     with open(source, "rb") as infile, open(target, "wb") as outfile:  
9  ✓         while buffer := infile.read(16 * MiB):  
10             outfile.write(buffer)
```



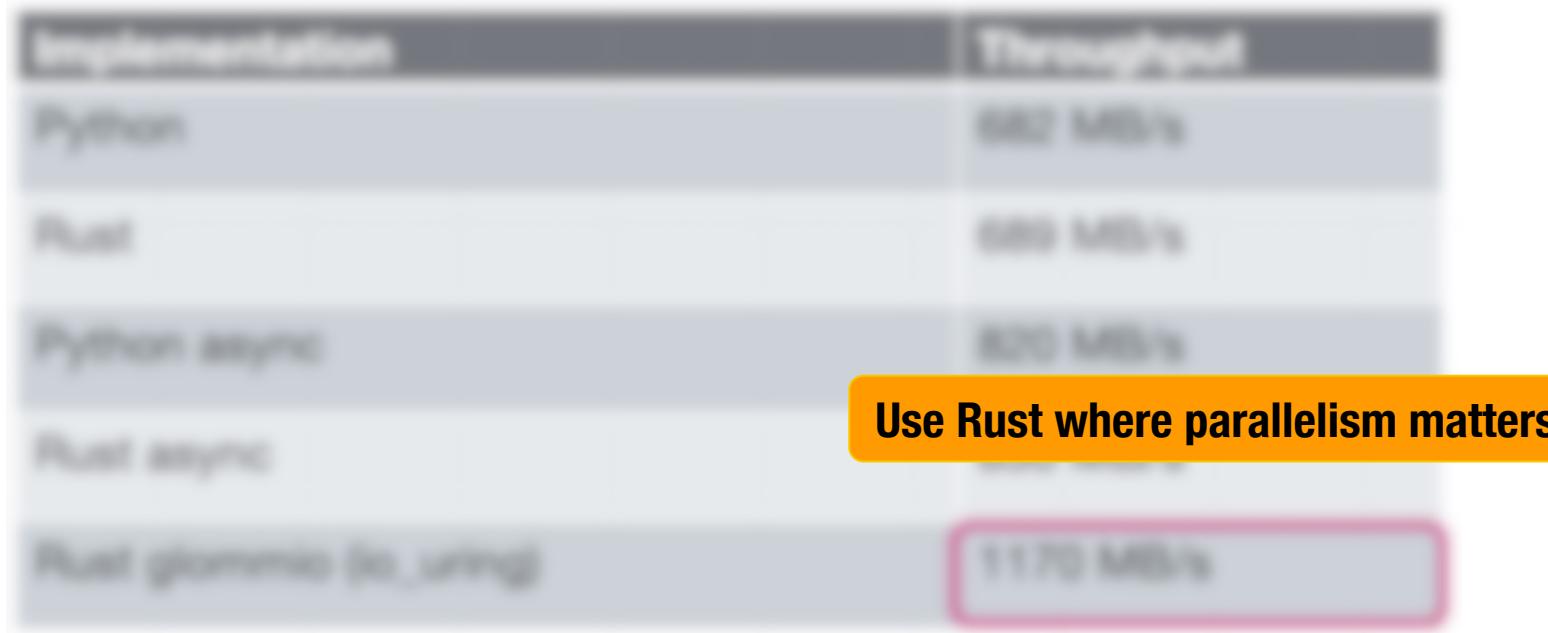
| Implementation | Throughput |
|----------------|------------|
| Python | 682 MB/s |
| Rust | 689 MB/s |

```
6  def main(source: str, target: str) -> None:  
7      """Copy source to target path"""  
8      with open(source, "rb") as infile, open(target, "wb") as outfile:  
9          while buffer := infile.read(16 * MiB):  
10              outfile.write(buffer)
```



| Implementation | Throughput |
|----------------|------------|
| Python | 682 MB/s |
| Rust | 689 MB/s |
| Python async | 820 MB/s |
| Rust async | 830 MB/s |

| Implementation | Throughput |
|-------------------------|------------|
| Python | 682 MB/s |
| Rust | 689 MB/s |
| Python async | 820 MB/s |
| Rust async | 830 MB/s |
| Rust glommio (io_uring) | 1170 MB/s |



**ONE DOES NOT SIMPLY
DEPLOY TO PRODUCTION**

Future proof your APIs by focussing on how they might break

ONE DOES NOT SIMPLY
DEPLOY TO PRODUCTION

Future proof your APIs by focussing on how they might break

Integration tests, integration tests and more integration tests

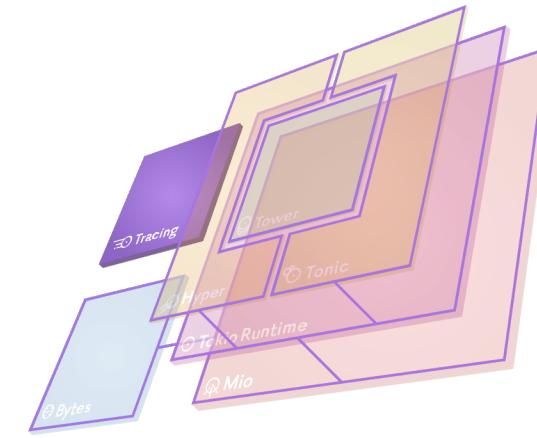
Future proof your APIs by focussing on how they might break

Integration tests, integration tests and more integration tests

Test as close to production as possible



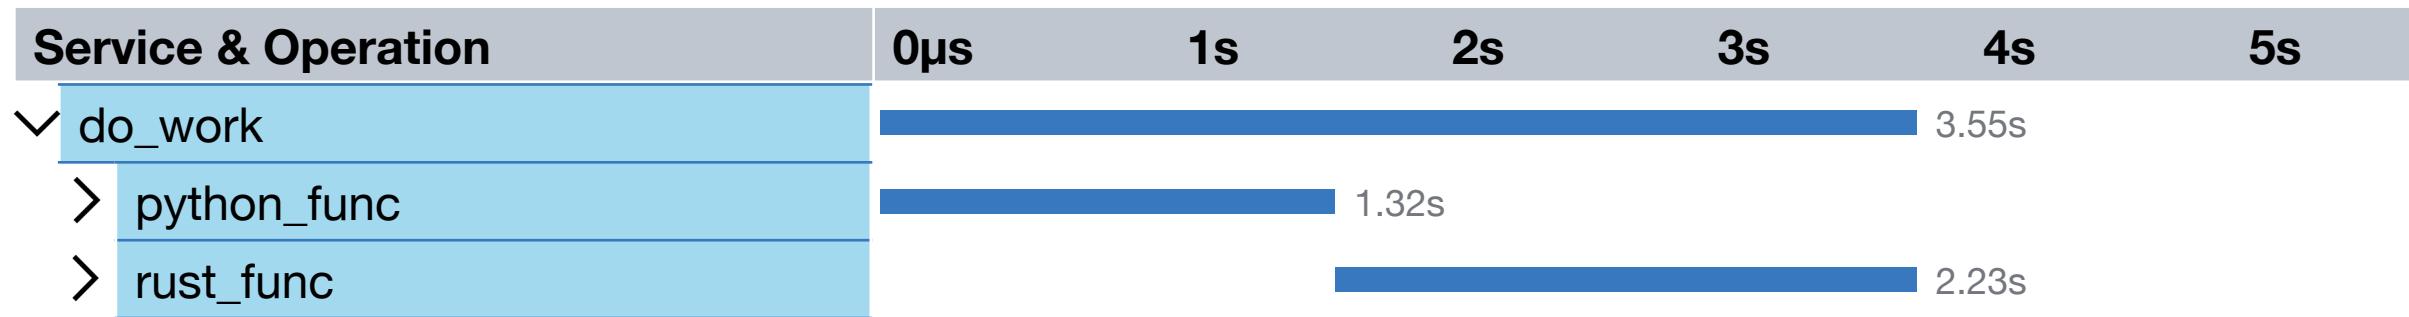
Structured, application-level diagnostics



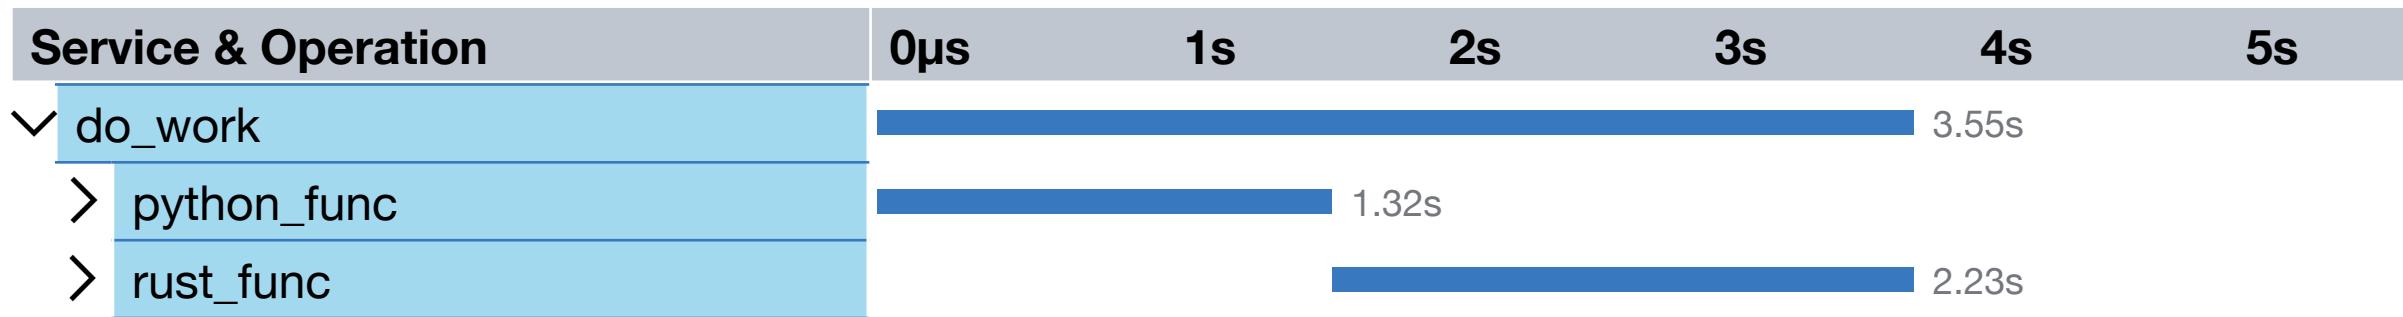
```
6   @tracer.start_as_current_span("do_work")
7   def do_work():
8       python_func()
9       rust_func()
```



```
6     @tracer.start_as_current_span("do_work")
7     def do_work():
8         python_func()
9         rust_func()
```



```
6     @tracer.start_as_current_span("do_work")
7     def do_work():
8         python_func()
9         rust_func(get_current_span().get_span_context().span_id)
```



```
32 #[pyfunction]
33 #[pyo3(name = "match_single_threaded")]
34 #[instrument(skip(py))]
35 pub fn main(py: Python, tasks: Vec<Task>, offers: Vec<Offer>) -> HashMap<String, String> {
```



```
32 #[pyfunction]
33 #[pyo3(name = "match_single_threaded")]
34 #[instrument(skip(py))]
35 pub fn main(py: Python, tasks: Vec<Task>, offers: Vec<Offer>) -> HashMap<String, String> {
36     set_python_span_as_parent(py);
37     py.allow_threads(move || match_tasks_to_offers(tasks, offers))
38 }
```



```
120  ✓ pub fn set_python_span_as_parent(py: Python) {  
121      let span: &Span = &tracing::Span::current();  
122      ✓ if span.is_disabled() {  
123          | return;  
124      }  
125  
126      let span_context: Result<Option<SpanContext>, ...> = get_python_opentelemetry_span_context(py);  
127  
128      ✓ if let Ok(Some(span_context: SpanContext)) = span_context {  
129          | let context: Context = opentelemetry::Context::new().with_remote_span_context(span_context);  
130          | span.set_parent(cx: context);  
131      }  
132  }  
133
```



```
120  ✓ pub fn set_python_span_as_parent(py: Python) {  
121      let span: &Span = &tracing::Span::current();  
122  ✓     if span.is_disabled() {  
123         return;  
124     }  
125  
126     let span_context: Result<Option<SpanContext>, ...> = get_python_opentelemetry_span_context(py);  
127  
128  ✓     if let Ok(Some(span_context: SpanContext)) = span_context {  
129         let context: Context = opentelemetry::Context::new().with_remote_span_context(span_context);  
130         span.set_parent(cx: context);  
131     }  
132 }  
133
```





```
63 pub fn get_python_opentelemetry_span_context(
64     py: Python,
65 ) -> PyResult<Option<opentelemetry::trace::SpanContext>> {
66     static OPENTELEMETRY_TRACE_IMPORT: OnceCell<Py<PyModule>> = OnceCell::new();
67
68     let result: Result<&Py<PyModule>, PyErr> = OPENTELEMETRY_TRACE_IMPORT OnceCell<Py<P...
69         .get_or_try_init(|| py.import(name: "opentelemetry.trace").map(op: Into::into));
70     let module: &Bound<'_, PyModule> = match result { ...
71     };
72
73
74     // opentelemetry always returns a context, but maybe with an invalid span ID = 0
75     let python_span_context: Bound<'_, PyAny> = module &Bound<'_, PyModule>
76         .call_method0(name: intern!(py, "get_current_span"))? Bound<'_, PyAny>
77         .call_method0(name: intern!(py, "get_span_context"))?;
78
79     let trace_id: u128 = python_span_context Bound<'_, PyAny>
80         .getattr(attr_name: intern!(py, "trace_id"))? Bound<'_, PyAny>
81         .extract::<u128>()?;
82
83     if trace_id == 0 {
84         return Ok(None);
85     }
86 }
```



```
63 pub fn get_python_opentelemetry_span_context(
64     py: Python,
65 ) -> PyResult<Option<opentelemetry::trace::SpanContext>> {
66     static OPENTELEMETRY_TRACE_IMPORT: OnceCell<Py<PyModule>> = OnceCell::new();
67
68     let result: Result<&Py<PyModule>, PyErr> = OPENTELEMETRY_TRACE_IMPORT OnceCell<Py<P...
69         .get_or_try_init(|| py.import(name: "opentelemetry.trace").map(op: Into::into));
70     > let module: &Bound<'_, PyModule> = match result { ...
71     };
72
73
74     // opentelemetry always returns a context, but maybe with an invalid span ID = 0
75     let python_span_context: Bound<'_, PyAny> = module &Bound<'_, PyModule>
76         .call_method0(name: intern!(py, "get_current_span"))? Bound<'_, PyAny>
77         .call_method0(name: intern!(py, "get_span_context"))?;
78
79
80     let trace_id: u128 = python_span_context Bound<'_, PyAny>
81         .getattr(attr_name: intern!(py, "trace_id"))? Bound<'_, PyAny>
82         .extract::<u128>()?;
83
84
85     if trace_id == 0 {
86         return Ok(None);
87     }
88 }
```



```
63 pub fn get_python_opentelemetry_span_context(
64     py: Python,
65 ) -> PyResult<Option<opentelemetry::trace::SpanContext>> {
66     static OPENTELEMETRY_TRACE_IMPORT: OnceCell<Py<PyModule>> = OnceCell::new();
67
68     let result: Result<&Py<PyModule>, PyErr> = OPENTELEMETRY_TRACE_IMPORT OnceCell<Py<P...
69         .get_or_try_init(|| py.import(name: "opentelemetry.trace").map(op: Into::into));
70     > let module: &Bound<'_, PyModule> = match result { ...
71     };
72
73
74     // opentelemetry always returns a context, but maybe with an invalid span ID = 0
75     let python_span_context: Bound<'_, PyAny> = module &Bound<'_, PyModule>
76         .call_method0(name: intern!(py, "get_current_span"))? Bound<'_, PyAny>
77         .call_method0(name: intern!(py, "get_span_context"))?;
78
79     let trace_id: u128 = python_span_context Bound<'_, PyAny>
80         .getattr(attr_name: intern!(py, "trace_id"))? Bound<'_, PyAny>
81         .extract::<u128>()?;
82
83
84     if trace_id == 0 {
85         return Ok(None);
86     }
87 }
```



```
88     if trace_id == 0 {
89         return Ok(None);
90     }
99
100    // create context
101    let span_context: SpanContext = opentelemetry::trace::SpanContext::new(
102        opentelemetry::trace::TraceId::from(trace_id),
103        opentelemetry::trace::SpanId::from(span_id),
104        opentelemetry::trace::TraceFlags::new(trace_flags),
105        is_remote: true, ←
106        trace_state: opentelemetry::trace::TraceState::NONE,
107    );
108
109    // Ensure span context is valid
110    > if !span_context.is_valid() { ...
115
116    Ok(Some(span_context))
117 } fn get_python_opentelemetry_span_context
```

```
120  ✓ pub fn set_python_span_as_parent(py: Python) {  
121      let span: &Span = &tracing::Span::current();  
122  ✓     if span.is_disabled() {  
123         return;  
124     }  
125  
126     let span_context: Result<Option<SpanContext>, ...> = get_python_opentelemetry_span_context(py);  
127  
128  ✓     if let Ok(Some(span_context: SpanContext)) = span_context {  
129         let context: Context = opentelemetry::Context::new().with_remote_span_context(span_context);  
130         span.set_parent(cx: context);  
131     }  
132 }  
133
```

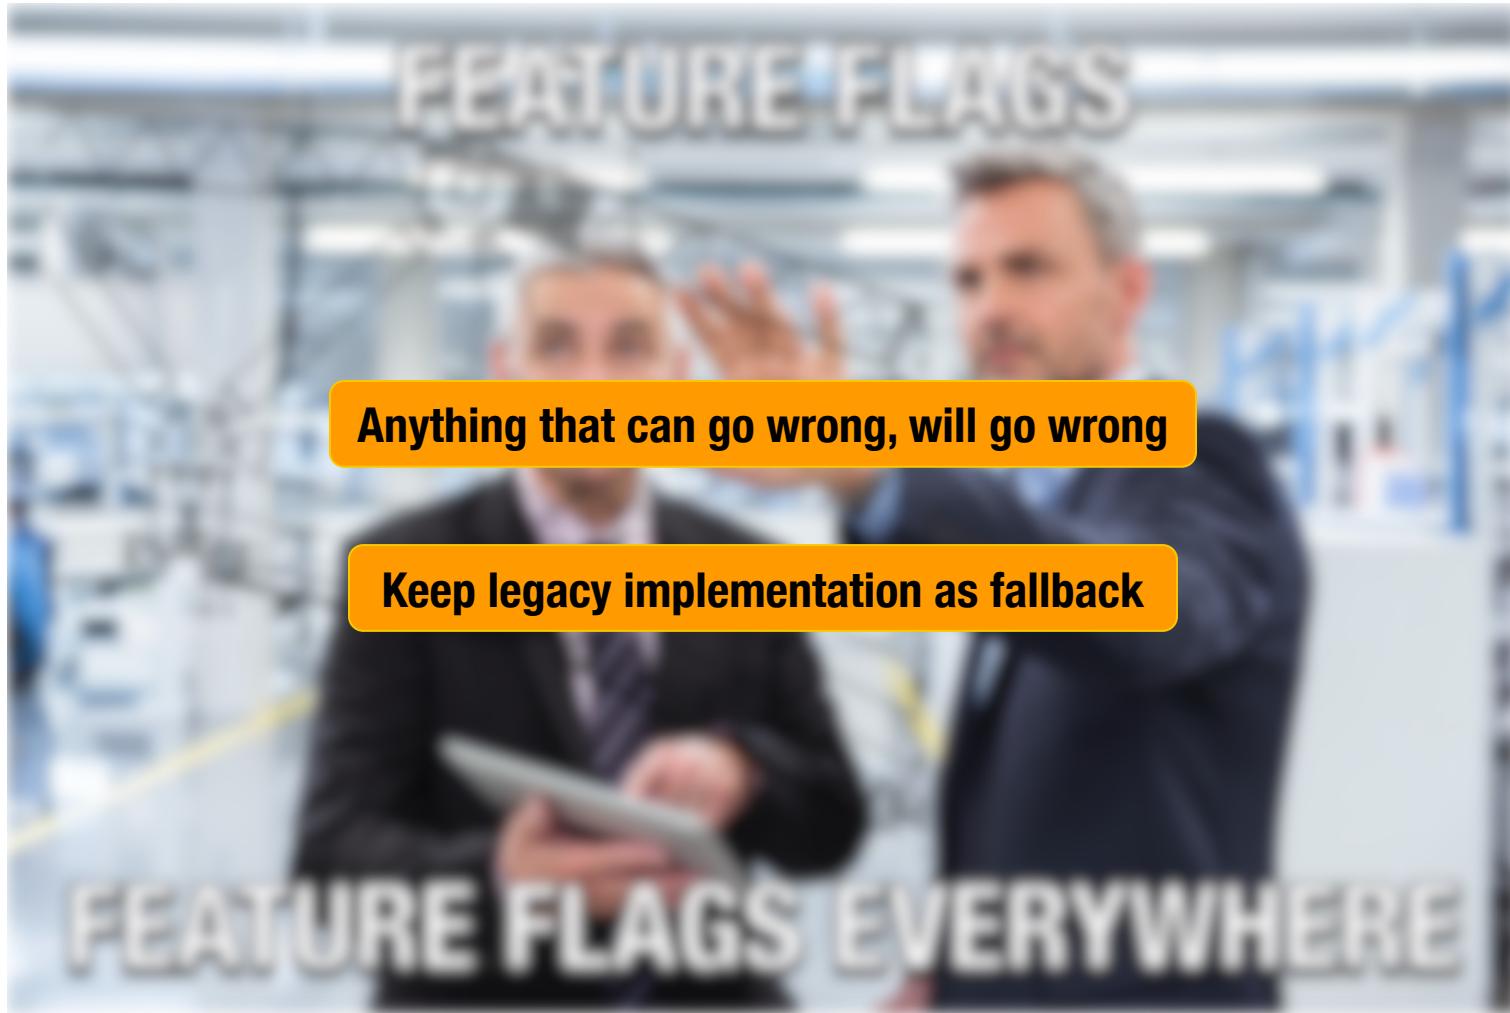


Build Rust using the release profile

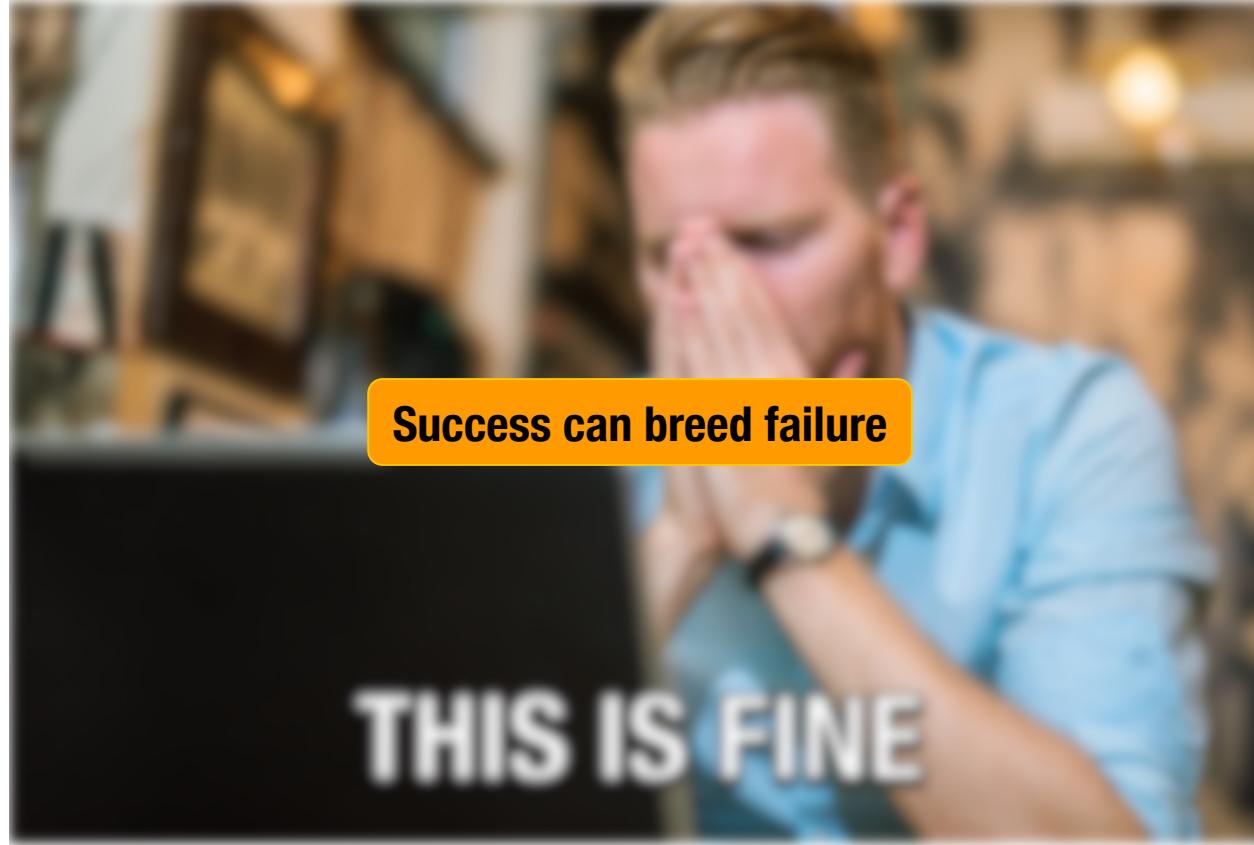
Version compatibility: PyO3 – rustc – python











2023
170K tasks/day
380K all time high



2025
340K tasks/day
978K all time high

Success can breed failure

Silver bullet

Focus on understanding and fixing the underlying root cause

Use Rust where parallelism matters

Release the GIL

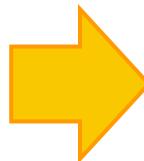
2023
170K tasks/day
380K all time high

Generate different test data (size, structure, etc.)

Scale matters – rust integration comes at a cost

Future proof your APIs by focussing on how they might break

Test as close to production as possible



Use existing toolset to its fullest potential

Anything that can go wrong, will go wrong

2025
340K tasks/day
978K all time high

Version compatibility: PyO3 – rustc – python

Keep legacy implementation as fallback

Integration tests, integration tests and more integration tests